


*COMPUTER VISION AND  
ITS APPLICATIONS  
CH 3\_EXERCISE*

徐繼聖

Gee-Sern Jison Hsu


National Taiwan University of Science  
and Technology





# Example 3-1: Convolution and Deconvolution

- Please download the “3-1\_Convolution and Deconvolution.zip” from the Moodle.
- Upload the “sample.jpg” to the Google Colab.
- Run the codes and get the output images through the convolution and deconvolution.



# Example 3-1: Convolution and Deconvolution

```
# First read the input image
img = cv2.imread('./sample.jpg')           Image path
img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
img = cv2.resize(img, (112, 112))         Resize the image to 112 by 112
plt.imshow(img)
img = torch.from_numpy(img.transpose((2, 0, 1)))
img_tensor = img.float().div(255).unsqueeze(0)
print('****Dimensions of input image****', img_tensor.shape)
```

We need to transpose the image in order to match the format of pytorch

The format of image:

Numpy: ( Height, Width, Channel)

Torch.Tensor( Channel, Height, Width)

# Example 3-1: Convolution and Deconvolution

Create the Network with 1 forward convolution

```
In [4]: #Define the neural Network
class Convolution(nn.Module):
    def __init__(self):
        super(Convolution, self).__init__()
        kernel = [[1/9, 1/9, 1/9],
                  [1/9, 1/9, 1/9],
                  [1/9, 1/9, 1/9]]
        kernel = torch.FloatTensor(kernel).unsqueeze(0).unsqueeze(0)
        self.weight = nn.Parameter(data=kernel, requires_grad=False)

    def forward(self, x):
        x1 = x[:, 0]
        x1 = F.conv2d(x1.unsqueeze(1), self.weight, stride=2, padding=0)

        return x1
```

Filter parameters

Single channel and create the forward path.

# Example 3-1: Convolution and Deconvolution

Feed the input image to the convolution

```
5): #create an instance of our FilterClass  
Convolution1 = Convolution()  
  
print('shape',img_tensor.shape)  
out = Convolution1(img_tensor)  
#out1=out.copy()  
print(out.size())  
#print(out)  
  
img_out = out.mul(255).byte()  
img_out = img_out.cpu().numpy().squeeze(0).transpose((1, 2, 0))  
print('*****Dimensions of output image*****:',img_out.shape)  
  
plt.imshow(img_out[:, :, 0])
```



Show the image.



# Example 3-1: Convolution and Deconvolution

Create the Deconvolution Network with 1 deconvolution layer

```
#Define the class Deconvolution
class Deconvolution(nn.Module):
    def __init__(self):
        super(Deconvolution, self).__init__()

        kernel = [[1/9, 1/9, 1/9],
                  [1/9, 1/9, 1/9],
                  [1/9, 1/9, 1/9]]
        #kernel3d = [kernel, kernel, kernel]


        kernel = torch.FloatTensor(kernel).unsqueeze(0).unsqueeze(0)
        self.weight = nn.Parameter(data=kernel, requires_grad=False)

    def forward(self, x):
        x1 = x[:, 0]
        x1 = F.conv_transpose2d(x1.unsqueeze(1), self.weight, stride=2, padding=0)
        return x1
```

# Example 3-1: Convolution and Deconvolution

Perform the deconvolution

```
#create an instance of our Deconvolution Class  
  
|  
Deconvolution1 = Deconvolution()  
out = Deconvolution1(out)  
print(out.size())  
#print(out)  
  
img_out = out.mul(255).byte()  
img_out = img_out.cpu().numpy().squeeze(0).transpose((1, 2, 0))  
print('*****Dimensions of output image*****:',img_out.shape)  
  
print(plt.imshow(img_out[:, :, 0]))
```



# Exercise 3-1: Convolution and Deconvolution

- Please download the “3-1\_Convolution and Deconvolution.zip” from the Moodle.
- Upload the “sample.jpg” to the Google Colab.
- Follow the example code and design a convolution kernel and a deconvolution kernel.
- Run the codes and get the output images through the designed convolution and deconvolution kernels.

Please write down your results and codes in MS Word, then upload to the Moodle.



## Exercise 3.2 : Autoencoder

- Please download the “exercise4.1\_Autoencoder.ipynb” on Moodle.
  - Train the autoencoder and compare the images that reconstruct from the different epoch.
  - Change the encoder and decoder to the below architecture and compare the difference.

Please copy your results and code and paste to a MS Word, then upload to Moodle.



# Example 3.2 : Autoencoder

Define the hyper-parameter and load the training data

```
num_epochs = 10  
batch_size = 32  
learning_rate = 1e-3
```

← Define the hyperparameter

```
img_transform = transforms.Compose([transforms.ToTensor()])
```

← Download the Mnist dataset to the folder './data'

```
dataset = torchvision.datasets.MNIST(root='./data', train=True, download=True, transform=img_transform)  
dataloader = torch.utils.data.DataLoader(dataset, batch_size=batch_size, shuffle=False)
```



# Example 3.2 : Autoencoder

Define the model architecture

```
class autoencoder(nn.Module):
    def __init__(self):
        super(autoencoder, self).__init__()
        self.encoder = nn.Sequential(
            nn.Linear(28 * 28, 128),
            nn.ReLU(True),
            nn.Linear(128, 64),
            nn.ReLU(True),
            nn.Linear(64, 12),
            nn.ReLU(True),
            nn.Linear(12, 3))
        self.decoder = nn.Sequential(
            nn.Linear(3, 12),
            nn.ReLU(True),
            nn.Linear(12, 64),
            nn.ReLU(True),
            nn.Linear(64, 128),
            nn.ReLU(True), nn.Linear(128, 28 * 28), nn.Tanh())
```

```
def forward(self, x):
    x = self.encoder(x)
    x = self.decoder(x)
    return x
```

```
autoencoder(
  (encoder): Sequential(
    (0): Linear(in_features=784, out_features=128, bias=True)
    (1): ReLU(inplace=True)
    (2): Linear(in_features=128, out_features=64, bias=True)
    (3): ReLU(inplace=True)
    (4): Linear(in_features=64, out_features=12, bias=True)
    (5): ReLU(inplace=True)
    (6): Linear(in_features=12, out_features=3, bias=True)
  )
  (decoder): Sequential(
    (0): Linear(in_features=3, out_features=12, bias=True)
    (1): ReLU(inplace=True)
    (2): Linear(in_features=12, out_features=64, bias=True)
    (3): ReLU(inplace=True)
    (4): Linear(in_features=64, out_features=128, bias=True)
    (5): ReLU(inplace=True)
    (6): Linear(in_features=128, out_features=784, bias=True)
    (7): Tanh()
  )
)
```

# Example 3.2 : Autoencoder

Define the model architecture

```
class autoencoder(nn.Module):  
    def __init__(self):  
        super(autoencoder, self).__init__()  
        self.encoder = nn.Sequential(  
            nn.Linear(28 * 28, 128),  
            nn.ReLU(True),  
            nn.Linear(128, 64),  
            nn.ReLU(True),  
            nn.Linear(64, 12),  
            nn.ReLU(True),  
            nn.Linear(12, 3))
```

Define the encoder

Define the decoder

```
        self.decoder = nn.Sequential(  
            nn.Linear(3, 12),  
            nn.ReLU(True),  
            nn.Linear(12, 64),  
            nn.ReLU(True),  
            nn.Linear(64, 128),  
            nn.ReLU(True),  
            nn.Linear(128, 28 * 28),  
            nn.Tanh())
```

```
    def forward(self, x):  
        x = self.encoder(x)  
        x = self.decoder(x)  
        return x
```



# Example 3.2 : Autoencoder

Define the model architecture

```
autoencoder(  
    (encoder): Sequential(  
      (0): Linear(in_features=784, out_features=128, bias=True)  
      (1): ReLU(inplace=True)  
      (2): Linear(in_features=128, out_features=64, bias=True)  
      (3): ReLU(inplace=True)  
      (4): Linear(in_features=64, out_features=12, bias=True)  
      (5): ReLU(inplace=True)  
      (6): Linear(in_features=12, out_features=3, bias=True)  
    )  
    (decoder): Sequential(  
      (0): Linear(in_features=3, out_features=12, bias=True)  
      (1): ReLU(inplace=True)  
      (2): Linear(in_features=12, out_features=64, bias=True)  
      (3): ReLU(inplace=True)  
      (4): Linear(in_features=64, out_features=128, bias=True)  
      (5): ReLU(inplace=True)  
      (6): Linear(in_features=128, out_features=784, bias=True)  
      (7): Tanh()  
    )  
  )  
)
```

# Example 3.2 : Autoencoder

→ Use the Mean Square Error as the loss function

```
criterion = nn.MSELoss()  
optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate, weight_decay=1e-5)
```

```
for epoch in range(num_epochs):  
    for data in dataloader:  
        img, _ = data  
        img = img.view(img.size(0), -1) → Flatten the input images  
        img = Variable(img).cuda()  
        output = model(img)  
        loss = criterion(output, img) → Compute the loss  
        optimizer.zero_grad()  
        loss.backward()  
        optimizer.step()  
    print('epoch [{ }/{ }], loss:{:.4f}'.format(epoch + 1, num_epochs, loss.item()))
```



# Example 3.2 : Autoencoder

```
if epoch % 1 == 0:
```

```
    img = to_img(img.cpu().data)
```

```
    save_image(img, './AE_img/input_{}.png'.format(epoch))
```


```
    pic = to_img(output.cpu().data)
```

```
    save_image(pic, './AE_img/output_{}.png'.format(epoch))
```

→ Save the input images

→ Save the reconstruction images

Input images



Epoch 1 : reconstruction images



Epoch 10 : reconstruction images



## Exercise 3.3 - VAE

- Please download the “exercise4.2\_VAE.ipynb” on Moodle.
- Please use different latent vector dimension
  1. 1 dimension
  2. 10 dimension
  3. 100 dimension

Please copy your results and code and paste to a MS Word, then upload to Moodle.



# Example 3.3 : VAE

Define the model architecture

Define the latent vector dimension

```
class VAE(nn.Module):  
    def __init__(self, image_channels=1, h_dim=1600, z_dim=20):  
        super(VAE, self).__init__()  
  
        self.conv1 = nn.Conv2d(image_channels, 32, kernel_size=4, stride=2)  
        self.conv2 = nn.Conv2d(32, 64, kernel_size=4, stride=2)  
  
        self.fc1 = nn.Linear(h_dim, z_dim)  
        self.fc2 = nn.Linear(h_dim, z_dim)  
        self.fc3 = nn.Linear(z_dim, h_dim)  
  
        self.deconv2 = nn.ConvTranspose2d(64, 32, kernel_size=5, stride=2)  
        self.deconv1 = nn.ConvTranspose2d(32, image_channels, kernel_size=4, stride=2)
```

# Example 3.3 : VAE

```
if epoch % 1 == 0:
```

```
    save = img.cpu().data
```

```
    save_image(img, './vae_img/input_{}.png'.format(epoch))
```

```
    save = recon_batch.cpu().data
```

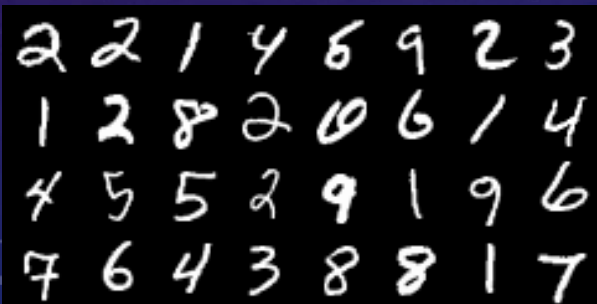
```
    save_image(save, './vae_img/output_{}.png'.format(epoch))
```

→ Save the input images

→ Save the reconstruction images


Epoch: 0

Input



2	2	1	4	6	9	2	3
1	2	8	2	0	6	1	4
4	5	5	2	9	1	9	6
7	6	4	3	8	8	1	7

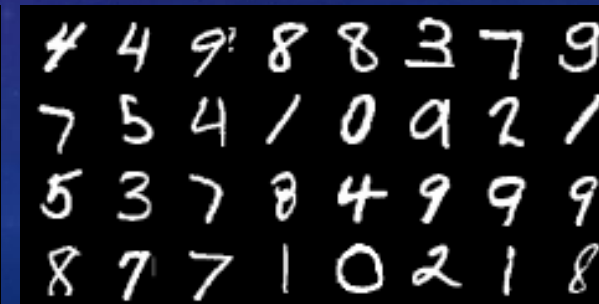
Reconstruct



0	2	1	4	6	9	2	3
1	2	8	2	0	6	1	4
4	5	5	2	9	1	9	6
7	6	4	3	8	8	1	7

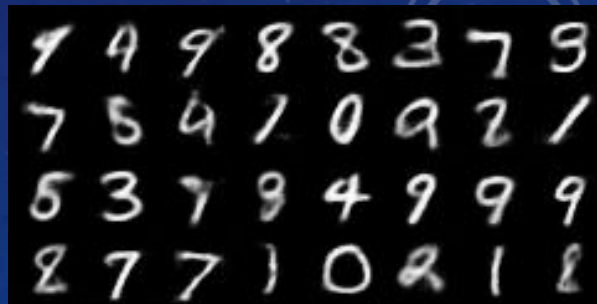
Epoch: 4

Input



4	4	9	8	8	3	7	9
7	5	4	1	0	9	2	1
5	3	7	8	4	9	9	9
8	7	7	1	0	2	1	8

Reconstruct



4	4	9	8	8	3	7	9
7	5	4	1	0	9	2	1
5	3	7	8	4	9	9	9
8	7	7	1	0	2	1	8