


*COMPUTER VISION AND  
ITS APPLICATION*

*CODING SKILLS  
FOR USING CNN*

徐繼聖

Gee-Sern Jison Hsu

National Taiwan University of Science  
and Technology





# Content Overview

- Feature Map Visualization and Its Dimension
- How To Training A Classifier and Codes of Building A Neural Network
- Other Examples For Studying
- CNN Review Problem


# Content Overview

- Feature Map Visualization and Its Dimension
- How To Training A Classifier and Codes of Building A Neural Network
- Other Examples For Studying
- CNN Review Problem

# Feature Map Visualization and Its Dimension

# Example 2-1: Feature Map Visualization

- Please download the “2-1\_Feature\_map\_visualization.zip” from the Moodle, which is built on the VGG-16 trained on the ImageNet.
- Upload the 2-1\_Feature\_map\_visualization.ipynb and imagenet1000\_clsidx\_to\_labels.txt to the Google Colab.
- Upload the “jellyfish.jpg” to the Google Colab.
- Run the codes and get the feature map from the selected layer.





# Example 2-1: Feature Map Visualization

imagenet1000\_clsidx\_to\_labels.txt


(in “Feature\_map\_visualization\_v2.7z”)

```
0: 'tench, Tinca tinca',
1: 'goldfish, Carassius auratus',
2: 'great white shark, white shark, man-eater, man-eating shark, Carcharodon carcharias',
3: 'tiger shark, Galeocerdo cuvieri',
4: 'hammerhead, hammerhead shark',
5: 'electric ray, crampfish, numbfish, torpedo',
6: 'stingray',
7: 'cock',
8: 'hen',
9: 'ostrich, Struthio camelus',
10: 'brambling, Fringilla montifringilla',
11: 'goldfinch, Carduelis carduelis',
12: 'house finch, linnet, Carpodacus mexicanus',
13: 'junco, snowbird',
14: 'indigo bunting, indigo finch, indigo bird, Passerina cyanea',
15: 'robin, American robin, Turdus migratorius',
16: 'bulbul',
17: 'jay',
18: 'magpie',
19: 'chickadee',
20: 'water ouzel, dipper',
21: 'kite',
22: 'bald eagle, American eagle, Haliaeetus leucocephalus',
23: 'vulture',
24: 'great grey owl, great gray owl, Strix nebulosa',
25: 'European fire salamander, Salamandra salamandra',
26: 'common newt, Triturus vulgaris',
27: 'eft',
28: 'spotted salamander, Ambystoma maculatum',
29: 'axolotl, mud puppy, Ambystoma mexicanum',
30: 'bullfrog, Rana catesbeiana',
31: 'tree frog, tree-frog',
32: 'tailed frog, bell toad, ribbed toad, tailed toad, Ascaphus trui',
33: 'loggerhead, loggerhead turtle, Caretta caretta',
```

# Example 2-1 : Feature Map Visualization



Original image: jellyfish.jpg



Probability of the classes

Probability TOP-3:

TOP\_1  
Probability: 0.9990069270133972  
Predicted: 'jellyfish'

TOP\_2  
Probability: 0.0008054533391259611  
Predicted: 'isopod'


TOP\_3  
Probability: 8.906585571821779e-05  
Predicted: 'chambered nautilus'

Predicted class : jellyfish

# Example 2-1 : Feature Map Visualization



Original image: jellyfish.jpg





# Example 2-1 : Coding Explanation

Overview of the sample code:

- Main Process for executing
- Function - FeatureVisualization
  - “\_\_init\_\_ (i.e. initialization)” for setting the pretrained model i.e., vgg16 on ImageNet
  - “process\_image” for the image preprocessing.
  - “get\_multi\_feature” for getting the feature maps.
  - “save\_feature\_to\_img” for saving the feature maps.
  - “Predict” for getting the prediction from the given image.

# Example 2-1 : Coding Explanation

## Main Process for executing

```
if __name__=='__main__':  
    # get class  
    c = {}  
    with open("imagenet1000_clsidx_to_labels.txt") as f:  
        for line in f:  
            (key, val) = line.split(":")  
            c[int(key)] = val.split(",")[0]  
    # Define image path and select the layer  
    myClass=FeatureVisualization('./jellyfish.jpg', 5)  
    print(myClass.pretrained_model2)  
    myClass.save_feature_to_img()  
    myClass.predict()
```

Open the txt file to get the information  
(Hint: You need to upload the file to Colab., or you will get the error "No such file.....")

Select the layer 5

Print the Network Architecture

1. Call the function to save the extracted feature from the selected layer in the VGG16.
2. Call the predict function to get the prediction from the given image



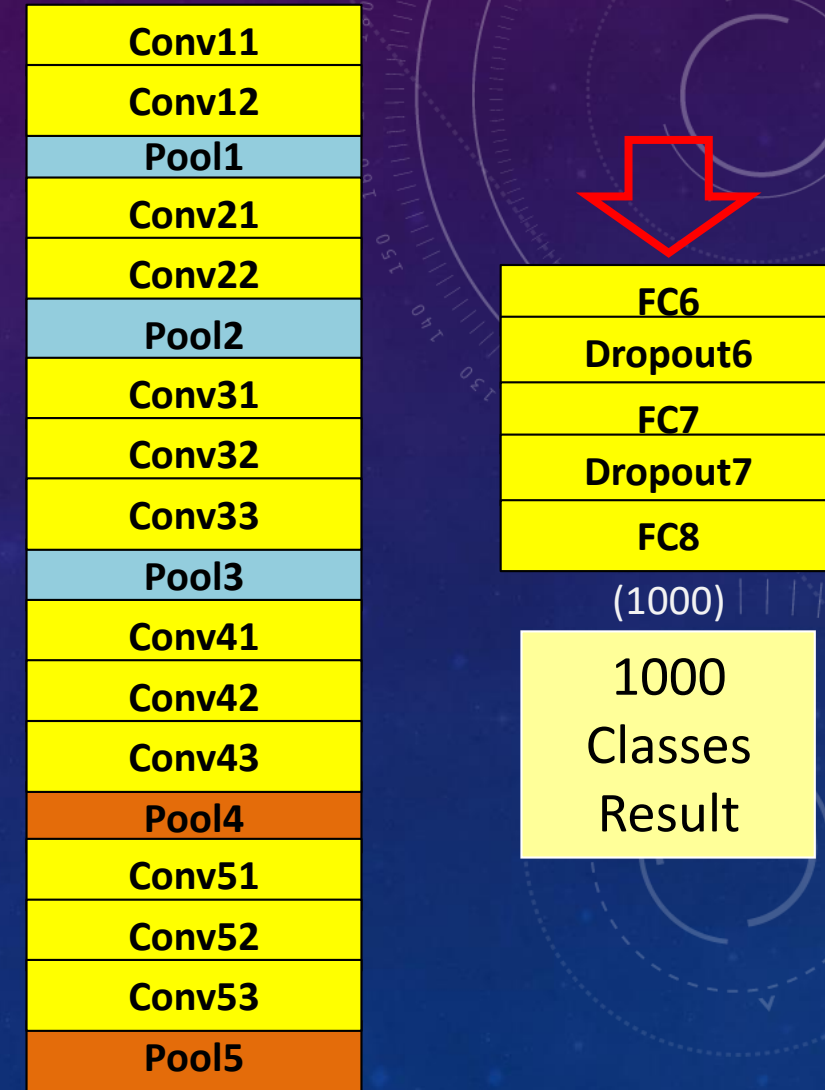
# Example 2-1 : Coding Explanation

## Feature Visualization

```
def __init__(self, img_path, selected_layer):  
    self.img_path=img_path  
    self.selected_layer=selected_layer  
    # Load pretrained model  
  
    self.pretrained_model = models.vgg16(pretrained=True).features  
    self.pretrained_model.eval()  
  
    self.pretrained_model2 = models.vgg16(pretrained=True)  
    self.pretrained_model2.eval()
```

Call the **feature part** of vgg16 pretrained model.  
“eval()” is for fixing the pretrained weight.

Call the **entire** vgg16 pretrained model (i.e. **the feature part and classifier part**)  
“eval()” is for fixing the pretrained weight.



Feature Extraction

# Example 2-1 : Coding Explanation

```
def process_image(self):  
    img=cv2.imread(self.img_path)  
    img=preprocess_image(img)  
    return img
```

Read the given image and preprocess it before feeding it into the model.

```
def preprocess_image(cv2im, resize_im=True):
```

```
    # Resize image  
    1. if resize_im:  
        cv2im = cv2.resize(cv2im, (224, 224))  
    im_as_arr = np.float32(cv2im)  
    2. im_as_arr = np.ascontiguousarray(im_as_arr[..., ::-1])  
    im_as_arr = im_as_arr.transpose(2, 0, 1) # Convert array to D,W,H  
    # Normalize the channels  
    3. for channel, _ in enumerate(im_as_arr):  
        im_as_arr[channel] /= 255  
    # Convert to float tensor  
    im_as_ten = torch.from_numpy(im_as_arr).float()  
    # Add one more channel to the beginning. Tensor shape = 1,3,224,224  
    4. im_as_ten.unsqueeze_(0)  
    # Convert to Pytorch variable  
    im_as_var = Variable(im_as_ten, requires_grad=True)  
    return im_as_var
```

Preprocess:

1. Resize to the 224x224 (i.e. **VGG16 input size**)
2. Convert the dimension to match the format of PyTorch.
3. Normalize the value of the data (**From 0 to 1, i.e., divide data by 255**)
4. Convert the data type to PyTorch tensor type



# Example 2-1 : Coding Explanation

```
def get_feature(self):  
    # Image preprocessing  
    input=self.process_image()  
    #print("input.shape: {}".format(input.shape))  
    x=input  
    for index,layer in enumerate(self.pretrained_model):  
        x=layer(x)  
        #print("x: {}".format(x.shape))  
        if (index == self.selected_layer):  
            return x
```

Feed the preprocessed image into the feature part of VGG16 model to extract the feature from the given layer.  
(i.e., the index of layer equal to the given value)

```
def get_multi_feature(self):  
    # Get the feature map  
    features=self.get_feature()  
    #print(features.shape)  
    result_path = './feat_first' + str(self.selected_layer)  
  
    if not os.path.exists(result_path):  
        os.makedirs(result_path)  
    print("On layer:{}, We can get the {} feature maps".format(self.selected_layer, features.shape[1]))  
    #print(features.shape[1])  
    for i in range(features.shape[1]):  
        feature=features[:,i,:]  
        feature=feature.view(feature.shape[1],feature.shape[2])  
        feature = feature.data.numpy()  
        feature = 1.0 / (1 + np.exp(-1 * feature))  
        feature = np.round(feature * 255)  
        save_name = result_path + '/' + str(i) + '.jpg'  
        cv2.imwrite(save_name, feature)
```

1. Create the folder to save the feature map.
2. Use for loop to save the extracted feature maps

# Example 2-1 : Coding Explanation

```
def save_feature_to_img(self):  
    #to numpy  
    feature=self.get_single_feature()  
    self.get_multi_feature()  
    feature=feature.data.numpy()  
  
    #use sigmod to [0,1]  
    # print(feature[0])  
    feature= 1.0/(1+np.exp(-1*feature))  
  
    # to [0,255]  
    feature=np.round(feature*255)  
    #print(self.selected_layer)  
    save_name = './feat_first' + str(self.selected_layer) + '.jpg'  
    cv2.imwrite(save_name, feature)
```

Call the function of extracting feature maps

Save the sample feature map



# Example 2-1 : Coding Explanation

```
def predict(self):  
    input=self.process_image()  
    outputs = self.pretrained_model2(input)  
  
    s = torch.nn.Softmax(dim=1)  
    result = s(outputs)  
    self.plot_probablity(result)  
  
    prob, predicted = result.sort(1,descending=True)  
    prob = prob.data.numpy()  
  
    predicted = predicted.data.numpy()  
  
    print("Probablity TOP-3:\n")  
    print("")  
    for i in range(3):  
        print("TOP_"+str(i+1))  
        print("Probablity: {}".format(prob[0][i]))  
        print("Predicted: {}\n".format(c[int(predicted[0][i])]))  
    return outputs
```

Call the preprocessed data, feed it into the entire vgg16 pretrained model, and get the output from the classifier.

Call the **Softmax** function to transform the output value to the **probability** and plot the figure.

Sort the predicted probability and show the first three value and its class in the ImageNet .

# Exercise 2-1: Feature Map Visualization

- Please download the “2-1\_Feature\_map\_visualization.zip” from the Moodle, which is built on the VGG-16 trained on the ImageNet.
- Upload the 2-1\_Feature\_map\_visualization.ipynb and imagenet1000\_clsidx\_to\_labels.txt to the Google Colab.
- Choose your own images from Internet.
- Compare the feature maps that extract from layer 5 and observe the size and dimension of the feature maps.

Please write down your results and codes in MS Word, then upload to the Moodle.



## Example 2-2 : Feature Map Visualization

- Please download the “2-2\_Feature\_map\_visualization.zip” on the Moodle and choose your own images from Internet.
- Upload the 2-2\_Feature\_map\_visualization.ipynb and imagenet1000\_clsidx\_to\_labels.txt to the Google Colab.
- Upload the given images “g1.jpg” and “g2.jpg” to the Google Colab.
- Run the codes and get the probabilities of these images.



## Example 2-3: Feature Comparison

- Please download the “2-2\_Feature\_map\_visualization.zip” on the Moodle
- Upload the 2-2\_Feature\_map\_visualization.ipynb and imagenet1000\_clsidx\_to\_labels.txt to the Google Colab.
- Upload the given images “g3.jpg” and “g4.jpg” to the Google Colab.
- Run the codes and get the comparisons.





# Example 2-2 & 2-3: Feature Map Visualization

```
if __name__ == '__main__':
    # get class
    c = {}
    with open("imagenet1000_clsidx_to_labels.txt") as f:
        for line in f:
            (key, val) = line.split(":")
            c[int(key)] = val.split(",")[0]

    # Define image path and select the layer
    myClass=FeatureVisualization('./dog6.jpg',12)
    Compare=FeatureVisualization('./dog9.jpg',12)
    print(myClass.pretrained_model2)

    myClass.save_feature_to_img()
    Compare.save_feature_to_img1()
    print("The first picture classification predict:")
    myClass_vector = myClass.predict()
    print("The second picture classification predict:")
    Compare_vector = Compare.predict()
    #Define cosine similarity
    cos= nn.CosineSimilarity(dim=1)
    #Define Euclidean distance
    euclidean_dist = torch.dist(myClass_vector, Compare_vector, p=2)
    cosine_dist = 1-cos(myClass_vector, Compare_vector)
    print(" Verification:")
    if cosine_dist < 0.6:
        print("They are the same!")
        print("Their cosine_distance: {}".format(cosine_dist))
    else:
        print("They are not the same!")
        print("Their cosine_distance: {}".format(cosine_dist))

    print("Their euclidean_dist: {}".format(euclidean_dist))
```

← Calculate the Euclidean distance between different pictures  
Calculate the Cosine distance between different pictures

← Define the threshold



# Example 2-2 & 2-3 : Feature Overview

## Results:

On layer:2, We can get the 64 feature maps  
The first picture classification predict:  
Probability TOP-3:

TOP\_1  
Probability:0.9882549047470093  
Predicted: 'jellyfish'

TOP\_2  
Probability:0.00702690239995718  
Predicted: 'isopod'

TOP\_3  
Probability:0.0019321587169542909  
Predicted: 'nematode'



The first picture

The second picture classification predict:  
Probability TOP-3:

TOP\_1  
Probability:0.2607852518558502  
Predicted: 'sports car'

TOP\_2  
Probability:0.20074793696403503  
Predicted: 'beach wagon'

TOP\_3  
Probability:0.13690434396266937  
Predicted: 'convertible'

Verification:  
They are not the same!  
Their cosine\_similarity:tensor([0.0470], grad\_fn=<DivBackward0>)  
Their euclidean\_dist:135.32333374023438



The second picture

## Exercise 2-2: Feature Map Visualization

- Please download the “2-1\_Feature\_map\_visualization.zip” on the Moodle and choose your own images from Internet.
- Upload the 2-1\_Feature\_map\_visualization.ipynb and imagenet1000\_clsidx\_to\_labels.txt to the Google Colab.
- Compare the probability of the images that contain multi classes and different variations (pose, occlusion, age).
- Please write down results and your codes in MS Word to the Moodle





## Exercise 2-3: Feature Comparison

- Please download the “2-3\_Feature\_map\_visualization.zip” on the Moodle
- Upload the 2-3\_Feature\_map\_visualization.ipynb and imagenet1000\_clsidx\_to\_labels.txt to the Google Colab.
- Please use the ResNet-50 pretrained model provided by Pytorch Package.
- Compare the similarity of the images that **contain two classes and different variations (pose, occlusion, age)**.
- Please write down result and your code in MS Words to the Moodle







# **Codes of Building A Neural Network and How To Training A Classifier**

# Example 2-4: Build A Classifier

Define a Convolutional Neural Network

Input size :  $3 * 32 * 32$

Input channel

Output channel

```
class Net(nn.Module):  
    def __init__(self):  
        super(Net, self).__init__()   
        self.conv1 = nn.Conv2d(3, 6, 5)   
        self.pool = nn.MaxPool2d(2, 2)   
        self.conv2 = nn.Conv2d(6, 16, 5)   
        self.fc1 = nn.Linear(16 * 5 * 5, 120)   
        self.fc2 = nn.Linear(120, 84)   
        self.fc3 = nn.Linear(84, 10)   
  
    def forward(self, x):  
        x = self.pool(F.relu(self.conv1(x)))  
        x = self.pool(F.relu(self.conv2(x)))  
        x = x.view(-1, 16 * 5 * 5)   
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = self.fc3(x)  
        return x
```

Kernel size

View function

Define forward path

Input image

$3 * 32 * 32$

Convolution

$6 * 28 * 28$

Max pooling

$6 * 14 * 14$

Convolution

$16 * 10 * 10$

Max pooling

$16 * 5 * 5$

$$32 - 5 + 1 = 28$$

$$28 \div 2 = 14$$

$$14 - 5 + 1 = 10$$

$$10 \div 2 = 5$$

# View function

View function is a tool to Reshape the tensor.

```
import torch  
a = torch.range(1, 16)
```

a = a tensor has 16 elements from 1 to 16

Try to reshape the tensor to 4 x 4

```
a = a.view(4, 4)
```

**What is the meaning of parameter -1?**

If there is any situation that you don't know how many rows you want but are sure of the number of columns, then you can specify this with a -1. (Note that you can extend this to tensors with more dimensions. Only one of the axis value can be -1).

This is a way of telling the library: "give me a tensor that has these many columns and you compute the appropriate number of rows that is necessary to make this happen".

Note : after the reshape the total number of elements need to remain the same



# Example 2-4: Build A Classifier

```
def imshow(img):  
    img = img / 2 + 0.5 # unnormalize  
    npimg = img.numpy()  
    plt.imshow(np.transpose(npimg, (1, 2, 0)))  
    plt.show()
```

```
def save_checkpoint(state, save_dir, filename='checkpoint.pth'):  
    save_name = save_dir + '_{}'.format(filename)  
    torch.save(state, save_name)
```

Set the filename of the weight file and use the function “torch.save” to save the weight file (Hint: “torch.load” can be used for loading the weight file).

Normalize a tensor image with mean and standard deviation. Given mean: (mean[1],...,mean[n]) and std: (std[1],...,std[n]) for n channels, this transform will normalize each channel of the input torch.\*Tensor i.e.,  $\text{output[channel]} = (\text{input[channel]} - \text{mean[channel]}) / \text{std[channel]}$

```
transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize(0.5, 0.5, 0.5), (0.5, 0.5, 0.5)])
```

**CIFAR10 dataset in Pytorch**

```
trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)  
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4, shuffle=True, num_workers=2)
```

**Setting the datasets path**

```
testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=transform)  
testloader = torch.utils.data.DataLoader(testset, batch_size=4, shuffle=False, num_workers=2)
```

```
classes = ('plane', 'car', 'bird', 'cat',  
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

**There are 10 classes in CIFAR10**

```
net = Net()
```

# Example 2-4: Build A Classifier

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

Select Cross-Entropy loss for classification

Using SGD optimizer, learning rate :0.001, momentum: 0.9

```
for epoch in range(3): # loop over the dataset multiple times
```

Setting the value of epoch (Hint: range(3) means 0 to 2)

```
    running_loss = 0.0
```

```
    for i, data in enumerate(trainloader, 0)
```

Use for loop to iter the batch data.

```
        # get the inputs; data is a list of [inputs, labels]
```

```
        inputs, labels = data
```

Each batch data contains the image data and the corresponding label.

```
        # zero the parameter gradients
```

```
        optimizer.zero_grad()
```

```
        # forward + backward + optimize
```

```
        outputs = net(inputs)
```

```
        loss = criterion(outputs, labels)
```

```
        loss.backward()
```

```
        optimizer.step()
```

1. Feed the input into the model and get the prediction.
2. Use the defined loss function to calculate the loss between the prediction and the label.
3. Use backward() to compute the gradient and use the optimizer.step() to update the weight

```
    # print statistics
```

```
    running_loss += loss.item()
```

```
    if i % 2000 == 1999: # print every 2000 mini-batches
```

```
        print("Epoch : {} steps : {} Training Loss : {}".format(epoch + 1, i + 1, running_loss / 2000) )
```

```
        running_loss = 0.0
```

Print the loss every 2000 steps

```
    save_checkpoint({'net':net.state_dict(), 'test_epoch{}'.format(epoch+1)})
```

Call the function to save the weight file.

```
print('Finished Training')
```

# Example 2-4: Build A Classifier

```
dataiter = iter(testloader)
images, labels = dataiter.next()
```

Use next() to get 1 batch data as a test sample.

```
outputs = net(images)
_, predicted = torch.max(outputs, 1)

print('Predicted: ', ' '.join('%5s' % classes[predicted[j]] for j in range(4)))
correct = 0
total = 0

with torch.no_grad():
    for data in testloader:
        images, labels = data
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy : %d %%' % (100 * correct / total))
```

1. Feed the input data into the model and get the value of CIFAR10's 10 classes.
2. Use torch.max to get the maximum score and its index

1. Use for loop to get the entire testing data
2. Feed the testing data and calculate the accuracy



# Image Classification

**CS231n: Convolutional Neural Networks for Visual Recognition**

<http://cs231n.github.io/classification/>

# Example - Training A Classifier

## 1. Loading and normalizing CIFAR10

Setting the datasets path

```
transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
```

```
trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)
```

```
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4, shuffle=True, num_workers=2)
```

```
testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=transform)
```

```
testloader = torch.utils.data.DataLoader(testset, batch_size=4, shuffle=False, num_workers=2)
```

```
classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

There are 10 classes in CIFAR10

# Example - Training A Classifier

## 3. Define a Loss function and optimizer

Select Cross-Entropy loss for classification

```
criterion = nn.CrossEntropyLoss()
```

```
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

Using SGD optimizer



# Example - Training A Classifier

## 4. Train the network and save the checkpoint

```
for epoch in range(10):  
    running_loss = 0.0  
    for i, data in enumerate(trainloader, 0):  
        inputs, labels = data  
  
        optimizer.zero_grad()  
  
        outputs = net(inputs)  
        loss = criterion(outputs, labels)  
        loss.backward()  
        optimizer.step()  
  
        running_loss += loss.item()  
        if i % 2000 == 1999:  
            print("Epoch : {} steps : {} Training Loss : {}".format(epoch + 1, i + 1, running_loss / 2000) )  
    running_loss = 0.0  
    save_checkpoint({'net':net.state_dict()}, 'test_epoch{}'.format(epoch+1))
```

Define the epoch

Backpropagation the loss to update the weights.

Save the checkpoint

# Model Training

```
outputs = model(inputs)
loss = criterion(outputs, targets)
# Backward and optimize
optimizer.zero_grad()
loss.backward()
optimizer.step()
```

- `optimizer.zero_grad()`: Because every time a variable is back propagated through, the gradient will be accumulated instead of being replaced.
- `loss.backward()`: Backward
- `optimizer.step()`: Parameters update based on the current gradient.

# Example - Training A Classifier

## 5. Test the network on the test data

```
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
print('Accuracy of the network on the 10000 test images: %d %%' % (100 * correct / total))
```

→ In the testing stage, the weights are fixed.

→ Check if predicted is the same as the labeled (G.T.)

Result:

```
Epoch : 10 steps : 8000 Training Loss : 0.8649811625033617
Epoch : 10 steps : 10000 Training Loss : 0.8769527244269848
Epoch : 10 steps : 12000 Training Loss : 0.8830881424993277
Finished Training
Accuracy : 61 %
```



# Example: Result Comparison on CIFAR10

## Use Pretrain model testing CIFAR10 Dataset

```
transform_test)
testloader = torch.utils.data.DataLoader(testset, batch_size=100, shuffle=False)

classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

# Choose Your Testing Model
print('==> Building model..')
net = VGG('VGG16')
#net = ResNet18()
#net = GoogLeNet()
```

→ Load pretrain architecture

```
# Load checkpoint.
print('==> Loading pretrained model from checkpoint..')
assert os.path.isdir('checkpoint'), 'Error: no checkpoint directory found!'
checkpoint = torch.load('./checkpoint/VGG16.pth') #checkpoint path
net.load_state_dict(checkpoint['net'])
```

→ \*\*Load pretrain model

## Use Your training model testing CIFAR10 Dataset

```
# your train model net
net= Net()
print(net)
net = net.to(device)
if device == 'cuda':
    net = torch.nn.DataParallel(net)
    cudnn.benchmark = True
```

→ \*\*Load your training architecture

```
# Load your taining checkpoint.
checkpoint = torch.load('./checkpoint/test_epoch3_checkpoint.pth')
net.module.load_state_dict(checkpoint['net'])
```

### VGG16

```
-----Start Testing-----
Loss: 0.745 | Acc: 75.590% (7559/10000)
```

### ResNet18

```
-----Start Testing-----
Loss: 0.727 | Acc: 76.210% (7621/10000)
```

### GoogleNet

```
-----Start Testing-----
Loss: 0.938 | Acc: 68.100% (6810/10000)
```

### Your training model result

```
-----Start Testing-----
Loss: 2.372 | Acc: 47.000% (4700/10000)
```


→ \*\*Load your training model

# Comparison of Deep Networks



# ResNet

- Since AlexNet, the state-of-the-art CNN architecture is going deeper and deeper.
- However, increasing network depth does not work by simply stacking layers together. Deep networks are hard to train because of the notorious vanishing gradient problem
- The core idea of ResNet is introducing a so-called “identity shortcut connection” that skips one or more layers, as shown in the following figure:



# Why ResNets Work



deeplearning.ai

## Case Studies

---

## Why ResNets work

[https://www.youtube.com/watch?v=RYth6EbBUqM&ab\\_channel=Deeplearning.ai](https://www.youtube.com/watch?v=RYth6EbBUqM&ab_channel=Deeplearning.ai)[9:12]

## Example 2-5: Comparison of Deep Networks

- Considering the following networks
  - $N_1$  has 10 conv layers and 1 Fc layer with Setup-1, please train  $N_1$  on the CIFAR-10 for 3 epochs;
  - $N_2$  has 10 conv layers with specified shortcut connections and 1 Fc layer with Setup-2, please train  $N_2$  on the CIFAR-10 for 3 epochs;



# Example 2-5: Comparison of Deep Networks

## Loading and normalizing CIFAR10

Setting the datasets path

```
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))]
)

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                       download=True, transform=transform)

trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
                                          shuffle=True, num_workers=8)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform)

testloader = torch.utils.data.DataLoader(testset, batch_size=4,
                                         shuffle=False, num_workers=8)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Normalize a tensor image with mean and standard deviation. Given mean: (mean[1],...,mean[n]) and std: (std[1],...,std[n]) for n channels, this transform will normalize each channel of the input torch.\*Tensor i.e.,  $\text{output}[\text{channel}] = (\text{input}[\text{channel}] - \text{mean}[\text{channel}]) / \text{std}[\text{channel}]$

There are 10 classes in CIFAR10

# Example 2-5: Comparison of Deep Networks

Setup-1 Define network(without shortcut)(10conv 1fc)

```
net = Net().to(device)
```

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Sequential(
            nn.Conv2d(3, 64, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(64),
            nn.ReLU(),)
        self.conv2 = nn.Sequential(
            nn.Conv2d(64, 64, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(64),
            nn.ReLU(),)
```

```
        self.conv6 = nn.Sequential(
            nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(128),
            nn.ReLU(),)
        self.conv7 = nn.Sequential(
            nn.Conv2d(128, 128, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(128),
            nn.ReLU(),)
```

```
        self.conv10 = nn.Sequential(
            nn.Conv2d(128, 128, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(128),
            nn.ReLU(),)
        self.fc1 = nn.Linear(131072, 10)
```

```
def forward(self, x):
    x = self.conv1(x)
    x = self.conv2(x)
    x = self.conv3(x)
    x = self.conv4(x)
    x = self.conv5(x)
    x = self.conv6(x)
    x = self.conv7(x)
    x = self.conv8(x)
    x = self.conv9(x)
    x = self.conv10(x)

    x = x.view(-1, 131072)
    x = F.relu(self.fc1(x))

    return x
```

Define forward path

# Example 2-5: Comparison of Deep Networks

## Setup-1 Print network(without shortcut)(10conv 1fc)

```
Net(  
  (conv1): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU()  
  )  
  (conv2): Sequential(  
    (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU()  
  )  
  (conv3): Sequential(  
    (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU()  
  )  
  (conv4): Sequential(  
    (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU()  
  )  
  (conv5): Sequential(  
    (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (2): ReLU()  
  )  
)
```

```
(conv6): Sequential(  
  (0): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
  (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
  (2): ReLU()  
)  
(conv7): Sequential(  
  (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
  (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
  (2): ReLU()  
)  
(conv8): Sequential(  
  (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
  (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
  (2): ReLU()  
)  
(conv9): Sequential(  
  (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
  (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
  (2): ReLU()  
)  
(conv10): Sequential(  
  (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
  (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
  (2): ReLU()  
)  
(fc1): Linear(in_features=131072, out_features=10, bias=True)
```



# Example 2-5: Deep Network Comparison

## Setup-2 Define network(with shortcut)(10conv 1fc)

```
net = ResNet18().to(device)
```

```
class ResidualBlock(nn.Module):
    def __init__(self, inchannel, outchannel, stride=1):
        super(ResidualBlock, self).__init__()
        self.left = nn.Sequential(
            nn.Conv2d(inchannel, outchannel, kernel_size=3, stride=stride, padding=1, bias=False),
            nn.BatchNorm2d(outchannel),
            nn.ReLU(inplace=True),
            nn.Conv2d(outchannel, outchannel, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(outchannel)
        )
        self.shortcut = nn.Sequential()
        if stride != 1 or inchannel != outchannel:
            self.shortcut = nn.Sequential(
                nn.Conv2d(inchannel, outchannel, kernel_size=1, stride=stride, bias=False),
                nn.BatchNorm2d(outchannel)
            )

    def forward(self, x):
        out = self.left(x)
        out += self.shortcut(x)
        out = F.relu(out)
        return out
```

Define shortcut

```
def ResNet18():
    return ResNet(ResidualBlock)
```

```
class ResNet(nn.Module):
    def __init__(self, ResidualBlock, num_classes=10):
        super(ResNet, self).__init__()
        self.inchannel = 64
        self.conv1 = nn.Sequential(
            nn.Conv2d(3, 64, kernel_size=3, stride=1, padding=1, bias=False),
            nn.BatchNorm2d(64),
            nn.ReLU(),
        )
        self.layer1 = self.make_layer(ResidualBlock, 64, 2, stride=1)
        self.layer2 = self.make_layer(ResidualBlock, 128, 2, stride=2)
        #self.layer3 = self.make_layer(ResidualBlock, 256, 2, stride=2)
        #self.layer4 = self.make_layer(ResidualBlock, 512, 2, stride=2)
        self.fc = nn.Linear(2048, num_classes)

    def make_layer(self, block, channels, num_blocks, stride):
        strides = [stride] + [1] * (num_blocks - 1)
        layers = []
        for stride in strides:
            layers.append(block(self.inchannel, channels, stride))
            self.inchannel = channels
        return nn.Sequential(*layers)

    def forward(self, x):
        out = self.conv1(x)
        out = self.layer1(out)
        out = self.layer2(out)
        #out = self.layer3(out)
        #out = self.layer4(out)
        out = F.avg_pool2d(out, 4)
        out = out.view(out.size(0), -1)
        out = self.fc(out)
        return out
```

Define forward path

# Example 2-5: Deep Network Comparison

## Setup-2 Print network(with shortcut)(10conv 1fc)

```
ResNet(
  (conv1): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU()
  )
  (layer1): Sequential(
    (0): ResidualBlock(
      (left): Sequential(
        (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
      (shortcut): Sequential()
    )
    (1): ResidualBlock(
      (left): Sequential(
        (0): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
        (4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
      (shortcut): Sequential()
    )
  )
)
```

```
(layer2): Sequential(
  (0): ResidualBlock(
    (left): Sequential(
      (0): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU(inplace=True)
      (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (4): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (shortcut): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (1): ResidualBlock(
    (left): Sequential(
      (0): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (2): ReLU(inplace=True)
      (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (4): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (shortcut): Sequential()
  )
)
(fc): Linear(in_features=2048, out_features=10, bias=True)
```

shortcut

# Example 2-5: Comparison of Deep Networks

Define a Loss function and optimizer

Select Cross-Entropy loss for classification

```
criterion = nn.CrossEntropyLoss()  
optimizer = optim.Adam(net.parameters(), lr=0.001)
```

Using Adam optimizer



# Example 2-5: Comparison of Deep Networks

Train the network and save the checkpoint

```
for epoch in range(3): # loop over the dataset multiple times
```

Define the epoch

```
    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)
```

```
        # zero the parameter gradients
        optimizer.zero_grad()
```

```
        # forward + backward + optimize
        outputs = net(inputs)
```

```
        loss = criterion(outputs, labels)
```

```
        loss.backward()
```

```
        optimizer.step()
```

Backpropagation the loss to update the weights.

```
        # print statistics
```

```
        running_loss += loss.item()
```

```
        if i % 2000 == 1999: # print every 2000 mini-batches
```

```
            print("Epoch : {} steps : {} Training Loss : {}".format(epoch + 1, i + 1, running_loss / 2000) )
```

```
            running_loss = 0.0
```

```
            save_checkpoint({'net':net.state_dict()}, 'test_epoch{}'.format(epoch+1))
```

Save the checkpoint

```
print('Finished Training')
```

# Example 2-5: Comparison of Deep Networks

Test the network on the test data

```
correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        outputs = net(images)

        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy : %d %%' % (100 * correct / total))
```

→ In the testing stage, the weights are fixed.

→ Check if predicted is the same as the labeled (G.T.)

# Example 2-5: Comparison of Deep Networks

## Result(with shortcut)

```
Epoch : 1 steps : 2000 Training Loss : 2.314868042707443
Epoch : 1 steps : 4000 Training Loss : 2.3025851249694824
Epoch : 1 steps : 6000 Training Loss : 2.3025851249694824
Epoch : 1 steps : 8000 Training Loss : 2.3025851249694824
Epoch : 1 steps : 10000 Training Loss : 2.3025851249694824
Epoch : 1 steps : 12000 Training Loss : 2.3025851249694824
Epoch : 2 steps : 2000 Training Loss : 2.3025851249694824
Epoch : 2 steps : 4000 Training Loss : 2.3025851249694824
Epoch : 2 steps : 6000 Training Loss : 2.3025851249694824
Epoch : 2 steps : 8000 Training Loss : 2.3025851249694824
Epoch : 2 steps : 10000 Training Loss : 2.3025851249694824
Epoch : 2 steps : 12000 Training Loss : 2.3025851249694824
Epoch : 3 steps : 2000 Training Loss : 2.3025851249694824
Epoch : 3 steps : 4000 Training Loss : 2.3025851249694824
Epoch : 3 steps : 6000 Training Loss : 2.3025851249694824
Epoch : 3 steps : 8000 Training Loss : 2.3025851249694824
Epoch : 3 steps : 10000 Training Loss : 2.3025851249694824
Epoch : 3 steps : 12000 Training Loss : 2.3025851249694824
```

Finished Training



GroundTruth: cat ship ship plane


Predicted: horse horse horse horse

Accuracy : 10 %

## Result(without shortcut)

```
Epoch : 1 steps : 2000 Training Loss : 1.8878891510665416
Epoch : 1 steps : 4000 Training Loss : 1.4781636514812708
Epoch : 1 steps : 6000 Training Loss : 1.2983864627033472
Epoch : 1 steps : 8000 Training Loss : 1.1348232387583703
Epoch : 1 steps : 10000 Training Loss : 1.0744273342452944
Epoch : 1 steps : 12000 Training Loss : 0.9819176591066644
Epoch : 2 steps : 2000 Training Loss : 0.8823310074708425
Epoch : 2 steps : 4000 Training Loss : 0.8428827857617289
Epoch : 2 steps : 6000 Training Loss : 0.8335133502772077
Epoch : 2 steps : 8000 Training Loss : 0.8076976113315905
Epoch : 2 steps : 10000 Training Loss : 0.7748700085091405
Epoch : 2 steps : 12000 Training Loss : 0.7552551930428016
Epoch : 3 steps : 2000 Training Loss : 0.6636658706957823
Epoch : 3 steps : 4000 Training Loss : 0.6618142743564677
Epoch : 3 steps : 6000 Training Loss : 0.6509611685594427
Epoch : 3 steps : 8000 Training Loss : 0.6368713211108697
Epoch : 3 steps : 10000 Training Loss : 0.650435102526215
Epoch : 3 steps : 12000 Training Loss : 0.6216317716715858
```

Finished Training



GroundTruth: cat ship ship plane

Predicted: dog ship car plane

Accuracy : 76 %



## Exercise 2-5: Comparison of Deep Networks

- Please download 2-5 Deep network comparsion.ipynb from moodle
- Modify the **Net**'s architecture and make its gradient not disappear (can be deleted layer, or change the components)

Upload your observations, comments and your code to Moodle in a docx.

# Comparison of Latent Vector



## Example 2-6: Comparison of Latent Vector

- Use VGG16-pretrain to predict the image and compare the similarity between two images, when there is only one class and two classes, and the difference the scores when the crop is in different places.




# Example 2-6 Comparson1


Conv 1-2




Conv 2-2




Conv 3-2



Conv 1-2



Conv 2-2










Conv 3-2



Cosine Similarity:  
0.7003



# Example 2-6 Comparson1




Cosine Similarity:  
0.7871






# Example 2-6 Comparson2




Conv 2-2




Conv 3-2



Conv 2-2



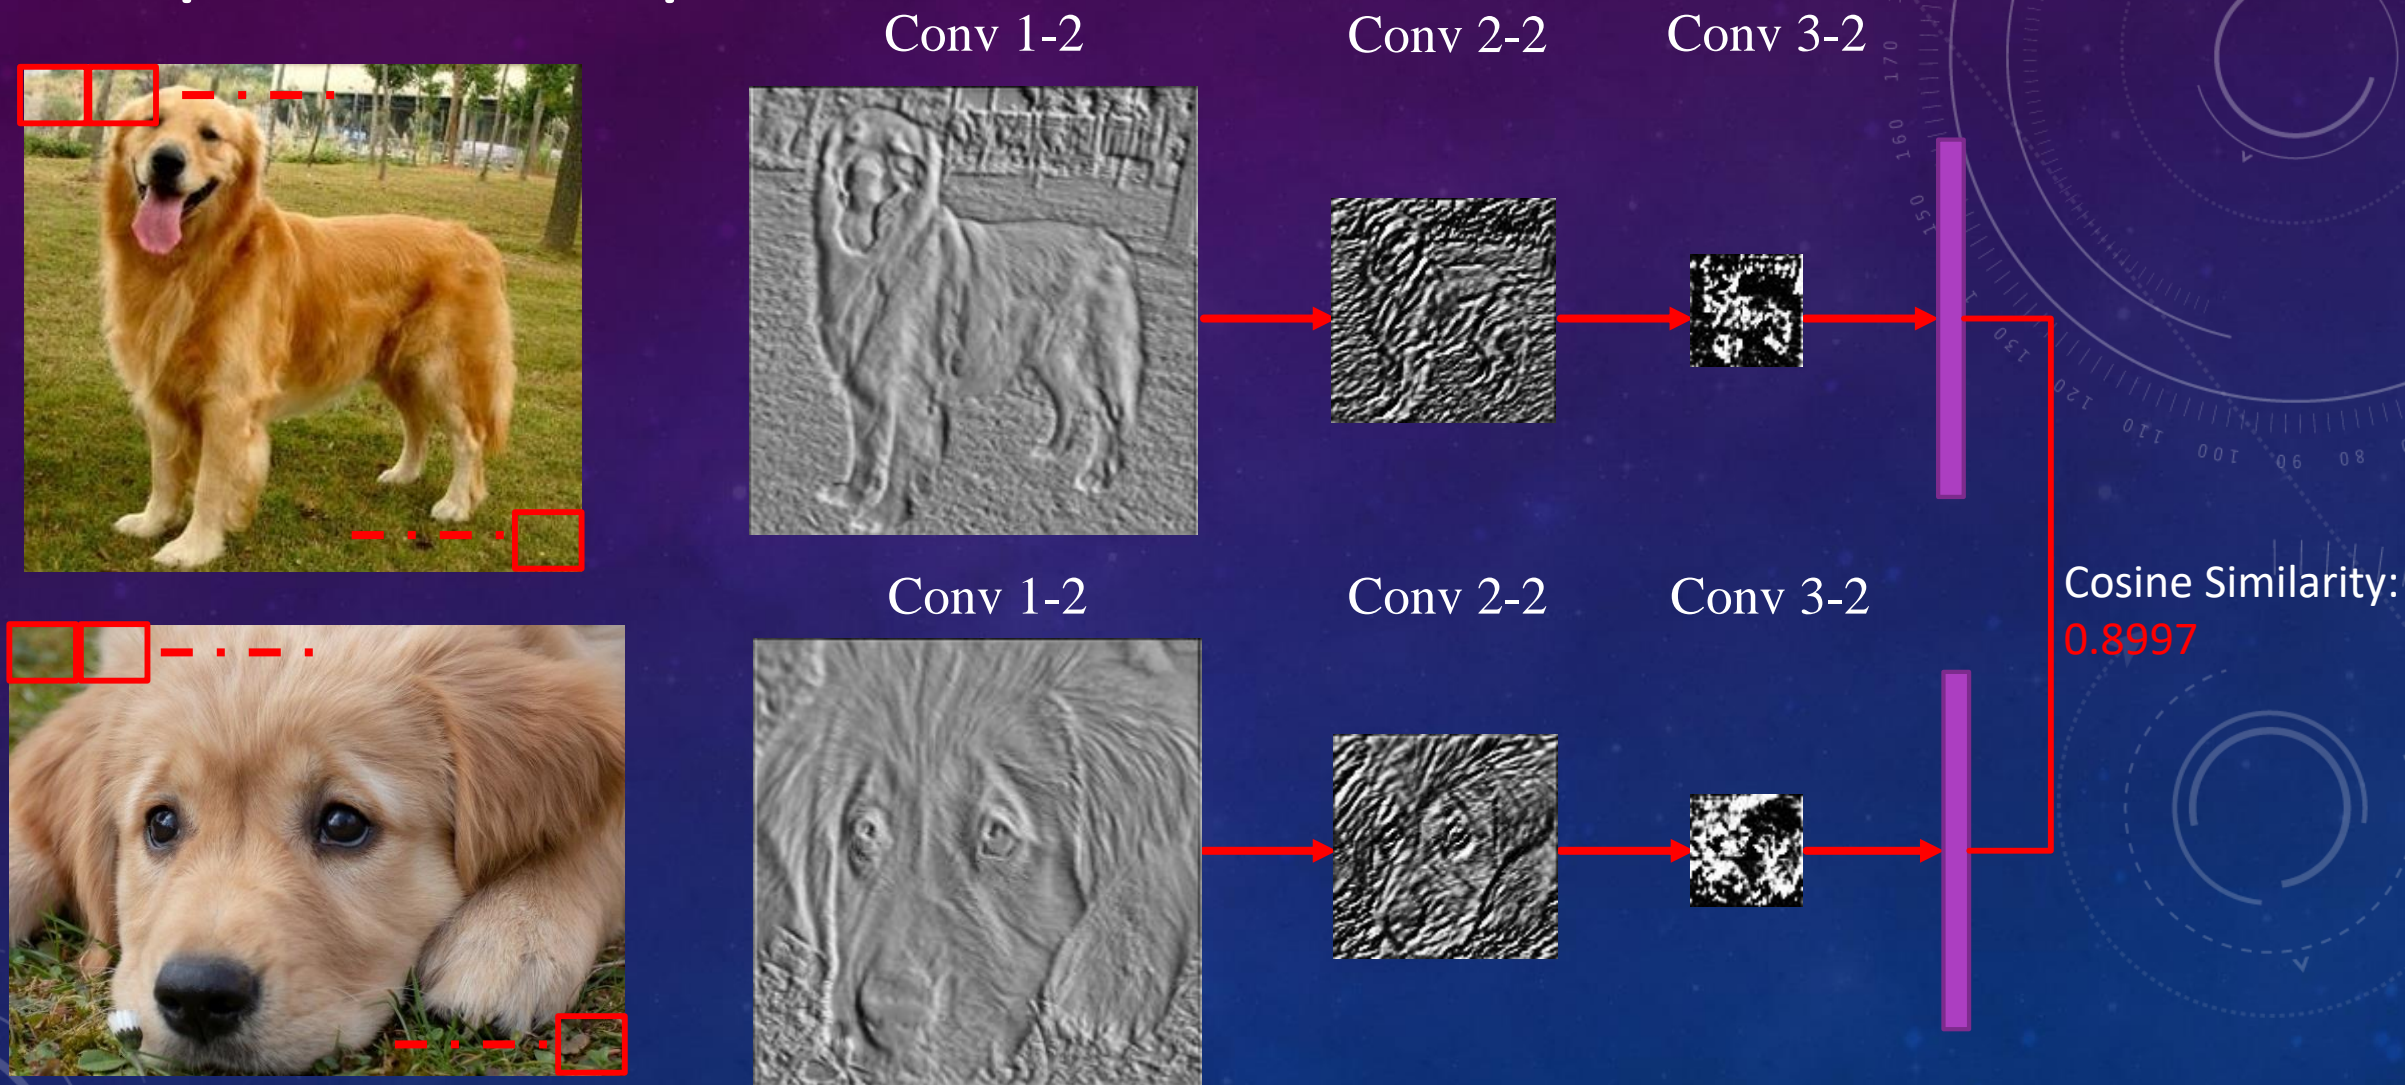
Conv 3-2





Cosine Similarity:  
0.8075



# Example 2-6 Comparson2




# Example 2-6 Comparson3




Conv 1-2

Conv 2-2




Conv 2-2


Conv 3-2




Conv 3-2



Conv 1-2



Conv 2-2








Conv 3-2



Cosine Similarity:  
**0.2748**









# Example 2-6 Comparson3



Cosine Similarity:  
0.7552



# Example 2-6 Comparson3



Cosine Similarity:  
0.8409

# Review Sample

# Problem 1 [30/100]

1. Prob1.ipynb gives you a VGG-16 trained on ImageNet. Upload the “imagenet1000\_clsidx\_to\_labels.txt “ and g1.jpg and g2.jpg to the Colab. Use Prob1.ipynb to show the following:
  - A. The feature maps and dimensions extracted from Layer 10. [8/30] (Example 2-1, Page 10)
  - B. Calculate the Euclidean distance between the images g1.jpg and g2.jpg, which are given with the code. [6/30] (Example 2-2 & 2-3, Page 19)
  - C. Please list the changes of dimension when feeding a image to the VGG-16. [16/30] (Example 2-1, Page 10)



# Problem 2 [45/100]

1. Please modify the Prob2.ipynb with the following requirements and set the Cross Entropy Loss, Adam Optimizer 0.002 learning rate and betas [0.5,0.999] to train a classifier :
  - A. Design a model with the following structure. (Example 2-4, Page 21)
    - First Conv. layer: Input: RGB, Output Channel 16, second Conv. layer: Output Channel 32, third Conv. layer: Output Channel 64.
    - FC-Layer1: Input: Defined by the third convolutional Layer, Output: 1200
    - FC-Layer2: Input: From FC- Layer1, Output: 600
    - FC-Layer3: Input: From FC- Layer2, Output: equal to your class size [16/40]
  - B. Save the model and name it as 'Prob2.pth' [4/40] (Example 2-4, Page 23)
  - C. Save the optimizer and name it as 'Prob2\_1.pth'[6/40] (Example 2-4, Page 23)

Net ( )

Layer type	Input channel	Output channel	Filter size	Stride	padding
Conv1	A	B	3	1	2
ReLU					
AvgPool			2	1	1
Conv2	B	C	2	1	1
ReLU					
AvgPool			2	2	1
Conv3	C	D	2	1	1
ReLU					
AvgPool			2	3	1
Linear1	E	G			
ELU					
Linear2	G	F			
ELU					
Linear3	F	G			

Please crop the parts that you modify in Prob2.ipynb and paste to the solution .docx.

# Problem 2

- E. Change the dataset to CIFAR10 and the learning rate :0.0002 [5/40] (Example 2-4, Page 26)
- F. Load the 'Prob2.pth and Prob2\_1.pth' obtained from C as pretrained model[8/40] (Hint: torch.load function, Page 33)
- G. Train the model on the CIFAR10 dataset[3/40] (Example 2-4, Page 26)
- H. Save the model and name it as 'Prob2\_2.pth' [3/40] (Example 2-4, Page 26)



## Problem 3 [25/100]

3. The output dimension of the feature map from Conv4 is  $64 \times 224 \times 256$ , please calculate the dimension of the Input and the feature maps from Conv1, Conv2, Conv3, Conv5

Layer type	Input channel	Output channel	Filter size	Stride
Conv1	3	8	3	1
AvgPool			4	1
Conv2	8	16	4	2
MaxPool			2	2
Conv3	16	32	2	1
MaxPool			2	1
Conv4	32	64	3	2
AvgPool			3	1
Conv5	64	128	7	1

Example 2-4, Page 24