



COMPUTER VISION AND ITS APPLICATION


Solutions-
Evaluation of Your Understanding on using Pytorch
for Setting up CNN

Jison Hsu
Artificial Vision Laboratory
Taiwan Tech

GRADES OVERVIEW

	Prob1	Prob2	Prob3	Total
B10630216	31	37	2	70
B10631030	25	25	15	65
B10603123	37	32	0	69
B10603043	15	32	15	62
D10803817	17	28	15	60
M10803816	18	29	7	54
M10815822	24	34	0	58
M10803339	22	28	15	65
D10903817	7	10	8	25
D10907801	21	39	8	68
M10903418	26	35	15	76
M10903429	31	0	15	46
M10903430	36	38	15	89
M10903807	21	37	4	62
M10903814	34	25	4	63
M10903151	16	14	2	32
M10903417	35	31	0	66
	24.47059	27.88235	8.235294	60.58824

Average score



Problem 1 [40/100]

1. Prob1.ipynb is a toy example, which shows you how to utilize the VGG-16 pre-trained model to extract the deep features of the input images. Please upload the files in “Prob1.zip” to the Google Colab, and answer the following questions:
 - 1) [2/40] Given the Q1.jpg as the input of the VGG-16, please show the feature maps and the feature dimensions of Layer-6 and Layer-10, and describe the differences.
 - 2) [5/40] Compute the Euclidean distance between the Q1.jpg and Q2.jpg using the Layer 6 feature maps.
 - 3) [5/40] Given the Q3.jpg as the input, please show the predicted class and top-3 probabilities.
 - 4) [2/40] Please compute the cosine similarities of the intra paired data (i.e., same classes): Q1.jpg and Q2.jpg
 - 5) [8/40] Please compute the cosine similarities of inter paired data (i.e., different classes):
 - 5.1) Q1.jpg and Q3.jpg
 - 5.2) Q2.jpg and Q4.jpg
 - 5.3) Q3.jpg and Q4.jpg

Problem 1 [40/100]

1. Prob1.ipynb is a toy example, which show you how to utilize the VGG-16 pre-trained model to extract the deep features of the input images. Please upload the files in “Prob1.zip” to the Google Colab., and answer the following questions:
 - 6) [6/40] Please describe what you observe in the four computed similarities from Prob 1.4) and 1.5).
 - 7) [6/40] Given a threshold = 0.5, please employ the computed similarities from Prob 1.4) and 1.5), and answer whether theses images are similar or not. (Hint: the similarity is greater than threshold, then they are similar !)
 - 8) [6/40] Given the predicted results in Prob1.7) and the ground-truth in Prob 1.4) and 1.5). Please show the accuracy rate.

Solution 1.1

Refer to
Example 2-1,
Problem 1.A

Given the Q1.jpg as the input of the VGG-16, please show the feature maps and the feature dimensions of Layer-6 and Layer-10, and describe the differences.



Q1.jpg

Solution 1.1

Feed the Q1.jpg into the VGG-16 pre-train model

```
Q1_6=FeatureVisualization('./Q1.jpg',6)  
Q1_10=FeatureVisualization('./Q1.jpg',10)
```

```
Q1_6.save_feature_to_img()  
Q1_10.save_feature_to_img1()
```

Selected Layer

Selected Layer

Save the extracted
feature maps


Observation:

It can be found that the details are not clear in the deeper layer, but you can still find the primary contour. It means that the model will distill and preserve the major features and dispel the unnecessary ones.


On layer:6, We can get the 128 feature maps
On layer:10, We can get the 256 feature maps



Q1.jpg



Layer6



Layer10

Solution 1.2

Refer to
[Example 2-1](#), [2-2](#),
[Problem 1.B](#)

Compute the Euclidean distance between the Q1.jpg and Q2.jpg using the Layer 6 feature maps.

```
Q2_6 = FeatureVisualization('./Q2.jpg',6)
A = Q1_6.get_feature()
B = Q2_6.get_feature()
Euclidean_dist = torch.dist(A,B,p=2)
print('Euclidean distance between the Q1.jpg and the Q2.jpg is {}'.format(Euclidean_dist))
```

```
def get_feature(self):
    # Image preprocessing
    input=self.process_image()
    #print("input.shape:{}".format(input.shape))
    x=input
    for index,layer in enumerate(self.pretrained_model):
        x=layer(x)
        # print("x:{}".format(x.shape))
        if (index == self.selected_layer):
            return x
```

Euclidean distance between the Q1.jpg and the Q2.jpg is 1017.1950073242188

Solution 1.2

Refer to
[Example 2-1](#), [2-2](#),
[Problem 1.B](#)

Feed the Q2.jpg into the VGG-16 pre-train model

```
Q2_6 = FeatureVisualization('./Q2.jpg', 6)
A = Q1_6.get_feature()
B = Q2_6.get_feature()
Euclidean_dist = torch.dist(A, B, p=2)
print('Euclidean distance between the Q1.jpg and the Q2.jpg is {}'.format(Euclidean_dist))
```

Q2.jpg and selected Layer6

Use get_feature() to
obtain the feature maps

Compute the Euclidean_distance

Euclidean distance between the Q1.jpg and the Q2.jpg is 1017.1950073242188

Solution 1.3

Refer to [Exercise 2-2](#)

Given the Q3.jpg as the input, please show the predicted class and top-3 probabilities.



Q3.jpg

Solution 1.3

Refer to [Exercise 2-2](#)

```
Q3 = FeatureVisualization('./Q3.jpg',0)
print("The first picture classification predict:")
Q3_predict = Q3.predict()
```

```
def predict(self):
    input=self.process_image()
    outputs = self.pretrained_model2(input)

    s = torch.nn.Softmax(dim=1)
    result = s(outputs)
    self.plot_probablity(result)

    prob, predicted = result.sort(1,descending=True)
    prob = prob.data.numpy()

    predicted = predicted.data.numpy()

    print("Probablity TOP-3:\n")
    print("")
    for i in range(3):

        print("TOP_"+str(i+1))
        print("Probablity: {}".format(prob[0][i]))
        print("Predicted: {}\n".format(c[int(predicted[0][i])]))

    return outputs
```



Q3.jpg

Probablity TOP-3:

TOP_1
Probablity:0.7808881402015686
Predicted: 'Egyptian cat'

TOP_2
Probablity:0.09954452514648438
Predicted: 'tabby'

TOP_3
Probablity:0.07643458992242813
Predicted: 'tiger cat'

Solution 1.3

[Refer to Exercise 2-2](#)

Feed the Q3.jpg into the VGG-16 pre-train model

```
Q3 = FeatureVisualization('./Q3.jpg',0)
print("The first picture classification predict:")
Q3_predict = Q3.predict()
```



Q3.jpg

Probability TOP-3:

```
TOP_1
Probability:0.7808881402015686
Predicted: 'Egyptian cat'

TOP_2
Probability:0.09954452514648438
Predicted: 'tabby'

TOP_3
Probability:0.07643458992242813
Predicted: 'tiger cat'
```

Call the predict() function and it will print the TOP-3 probabilities and their corresponding classes

Solution 1.3

Refer to [Exercise 2-2](#)

```
def predict(self):
    input=self.process_image()
    outputs = self.pretrained_model2(input)

    s = torch.nn.Softmax(dim=1)
    result = s(outputs)
    self.plot_probablity(result)

    prob, predicted = result.sort(1,descending=True)
    prob = prob.data.numpy()

    predicted = predicted.data.numpy()

    print("Probablity TOP-3:\n")
    print("")
    for i in range(3):

        print("TOP_"+str(i+1))
        print("Probablity: {}".format(prob[0][i]))
        print("Predicted: {}\n".format(c[int(predicted[0][i])]))
    return outputs
```

Get the outputs and transfer them to the probabilities by SoftMax function.

Sort these probabilities.

Show the top-3 probabilities and their corresponding classes


Solution 1.4

Refer to
Example 2-3
Exercise 2-3

Please compute the cosine similarities of the intra paired data (i.e., same classes): Q1.jpg and Q2.jpg



Q1.jpg



Q2.jpg

Intra Pair data

Solution 1.4

Refer to
[Example 2-3](#)
[Exercise 2-3](#)

Extract the latent feature by using `extract_latent_feature_vgg()`

```
Q1_latent = Q1_6.extract_latent_feature_vgg()
Q2_latent = Q2_6.extract_latent_feature_vgg()
cos = nn.CosineSimilarity(dim=1)

Q12_cosim = cos(Q1_latent, Q2_latent)
print('The cosine similarity between Q1 and Q2 is {}'.format(Q12_cosim.data))
```

Compute and show the similarity between Q1.jpg and Q2.jpg.

```
def extract_latent_feature_vgg(self):
    # Image preprocessing
    input=self.process_image()
    x=input


    x = self.pretrained_model2.features(x)
    x = self.pretrained_model2.avgpool(x)
    x = torch.flatten(x, 1)

    for index, layer in enumerate(self.pretrained_model2.classifier):
        x=layer(x)
        if (index == len(self.pretrained_model2.classifier)-2):
            return x
```

Extract the last second FC layer as the latent feature.



Q1.jpg



Q2.jpg

Intra Pair data

The cosine similarity between Q1 and Q2 is tensor([0.5561])

Solution 1.4

Refer to
[Example 2-3](#)
[Exercise 2-3](#)

Extract the latent feature by using `extract_latent_feature_vgg()`

```
Q1_latent = Q1_6.extract_latent_feature_vgg()  
Q2_latent = Q2_6.extract_latent_feature_vgg()  
cos = nn.CosineSimilarity(dim=1)
```

Define the cosine
similarity function

```
Q12_cosim = cos(Q1_latent, Q2_latent)  
print('The cosine similarity between Q1 and Q2 is {}'.format(Q12_cosim.data))
```

Compute and show the similarity between Q1.jpg and Q2.jpg.

Solution 1.4

Refer to
[Example 2-3](#)
[Exercise 2-3](#)

```
def extract_latent_feature_vgg(self):  
    # Image preprocessing  
    input=self.process_image()  
    x=input  
  
    x = self.pretrained_model2.features(x)  
    x = self.pretrained_model2.avgpool(x)  
    x = torch.flatten(x, 1)  
  
    for index, layer in enumerate(self.pretrained_model2.classifier):  
        x=layer(x)  
        if (index == len(self.pretrained_model2.classifier)-2):  
            return x
```

Extract the last second FC layer as the latent feature.

Solution 1.5

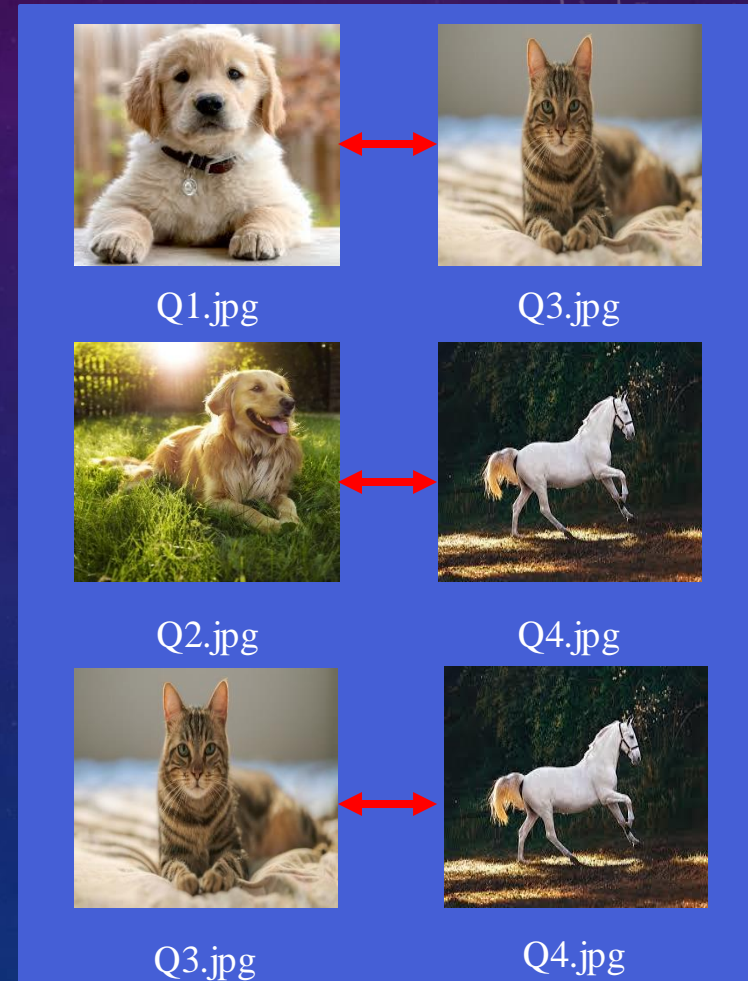
Refer to
Example 2-3
Exercise 2-3

Please compute the cosine similarities of inter paired data
(i.e., different classes):

5.1) Q1.jpg and Q3.jpg

5.2) Q2.jpg and Q4.jpg

5.3) Q3.jpg and Q4.jpg



Inter Pair data

Solution 1.5

Refer to
Example 2-3
Exercise 2-3

```
Q4 = FeatureVisualization('./Q4.jpg',0)
Q3_latent = Q3.extract_latent_feature_vgg()
Q4_latent = Q4.extract_latent_feature_vgg()

Q13_cosim = cos(Q1_latent,Q3_latent)
print('The cosine similarity between Q1 and Q3 is {}'.format(Q13_cosim.data))

Q24_cosim = cos(Q2_latent,Q4_latent)
print('The cosine similarity between Q2 and Q4 is {}'.format(Q24_cosim.data))

Q34_cosim = cos(Q3_latent,Q4_latent)
print('The cosine similarity between Q3 and Q4 is {}'.format(Q34_cosim.data))
```



```
The cosine similarity between Q1 and Q3 is tensor([0.0907])
The cosine similarity between Q2 and Q4 is tensor([0.3942])
The cosine similarity between Q3 and Q4 is tensor([0.1344])
```



Q1.jpg



Q3.jpg



Q2.jpg



Q4.jpg



Q3.jpg



Q4.jpg



Inter Pair data

Solution 1.6

Please describe what you observe in the four computed similarities from Prob 1.4) and 1.5).

```
The cosine similarity between Q1 and Q2 is tensor([0.5561])
```

```
The cosine similarity between Q1 and Q3 is tensor([0.0907])  
The cosine similarity between Q2 and Q4 is tensor([0.3942])  
The cosine similarity between Q3 and Q4 is tensor([0.1344])
```


Observation:

It can be seen that the similarity will be lower when comparing the different species. The value would be affected by the compositions of the two images.

Refer to
Example 2-3
Exercise 2-3



Q1.jpg



Q2.jpg



Q3.jpg



Q4.jpg

Solution 1.7

Refer to
Example 2-3
Exercise 2-3

Given a threshold = 0.5, please employ the computed similarities from Prob 1.4) and 1.5), and answer whether these images are similar or not. (Hint: the similarity is greater than threshold, then they are similar !)

Solution 1.7

Refer to
Example 2-3
Exercise 2-3

```
threshold = 0.5
## Q1.jpg and Q2.jpg
if Q12_cosim >= threshold:
    print("Q1.jpg and Q2.jpg are similar, its similarity is {}".format(Q12_cosim.data))
else :
    print("Q1.jpg and Q3.jpg are not similar, its similarity is {}".format(Q12_cosim.data))
## Q1.jpg and Q3.jpg
if Q13_cosim >= threshold:
    print("Q1.jpg and Q3.jpg are similar, its similarity is {}".format(Q13_cosim.data))
else :
    print("Q1.jpg and Q3.jpg are not similar, its similarity is {}".format(Q13_cosim.data))
## Q2.jpg and Q4.jpg
if Q24_cosim >= threshold:
    print("Q2.jpg and Q4.jpg are similar, its similarity is {}".format(Q24_cosim.data))
else :
    print("Q2.jpg and Q4.jpg are not similar, its similarity is {}".format(Q24_cosim.data))
## Q3.jpg and Q4.jpg
if Q34_cosim >= threshold:
    print("Q3.jpg and Q4.jpg are similar, its similarity is {}".format(Q34_cosim.data))
else :
    print("Q3.jpg and Q4.jpg are not similar, its similarity is {}".format(Q34_cosim.data))
```

Q1.jpg and Q2.jpg are similar, its similarity is tensor([0.5561])
Q1.jpg and Q3.jpg are not similar, its similarity is tensor([0.0907])
Q2.jpg and Q4.jpg are not similar, its similarity is tensor([0.3942])
Q3.jpg and Q4.jpg are not similar, its similarity is tensor([0.1344])

Solution 1.8

Refer to
[Example 2-3](#)
[Exercise 2-3](#)

Given the predicted results in Prob1.7) and the ground-truth in Prob 1.4) and 1.5). Please show the accuracy rate.

```
threshold= 0.5
similarities = [Q12_cosim,Q13_cosim,Q24_cosim,Q34_cosim]
pair_Groundtruth = [1,0,0,0]
correct = 0
Total_pair = 4
for sim,GT in zip(similarities,pair_Groundtruth):
    if sim >= threshold :
        pred = 1
    else :
        pred = 0
    if pred == GT:
        correct +=1
Accuracy = correct/Total_pair
print('The accuracy rate is {} %'.format(Accuracy*100))
```

The accuracy rate is 100.0 %

Solution 1.8

Refer to
Example 2-3
Exercise 2-3

```
threshold= 0.5
similarities = [Q12_cosim,Q13_cosim,Q24_cosim,Q34_cosim]
pair_Groundtruth = [1,0,0,0]
correct = 0
Total_pair = 4
for sim,GT in zip(similarities,pair_Groundtruth):
    if sim >= threshold :
        pred = 1
    else :
        pred = 0
    if pred == GT:
        correct +=1
Accuracy = correct/Total_pair
print('The accuracy rate is {} %'.format(Accuracy*100))
```

The similarities are obtained
from 1.4 and 1.5

GT of the paired data,
intra = 1, inter = 0

Compare the similarities
with threshold

Check whether the
prediction is equal to the GT

The accuracy rate is 100.0 %

Problem 2 [45/100]

2. Prob2.ipynb shows you how to train a classifier. Please modify the Prob2.ipynb with the following requirements and set the Cross Entropy Loss, optimizer SGD, 0.002 learning rate and 0.9 momentum to train a Fashion-MNIST classifier :
 - 1) [8/45] Design a model with the following structure and complete the table in the next page.
 - First Conv. layer: Input: Gray scale, Output Channel 32, second Conv. layer: Output Channel 64, third Conv. layer: Output Channel 128.
 - FC-Layer1: Input: Defined by the third convolutional Layer, Output: 1000
 - FC-Layer2: Input: From FC- Layer1, Output: 500
 - FC-Layer3: Input: From FC- Layer2, Output: equal to your class size
 - 2) [2/45] Train the classifier and show the losses during the training process.
 - 3) [2/45] Save the model and name it as 'Prob2.pth'.
 - 4) [2/45] Save the optimizer and name it as 'Prob2_1.pth'.

Net ()

Layer type	Input channel	Output channel	Filter size	Stride	padding
Conv1	A	B	3	1	1
ReLU					
MaxPool			2	2	0
Conv2	B	C	3	2	0
ReLU					
MaxPool			2	2	0
Conv3	C	D	4	2	1
ReLU					
AvgPool			2	1	1
Linear1	E	G			
ELU					
Linear2	G	F			
ELU					
Linear3	F	G			

Please crop the parts that you modify in Prob2.ipynb and paste to the solution.docx.

Problem 2

2. Prob2.ipynb shows how to train a classifier. Please modify the Prob2.ipynb with the following requirements and set the Cross Entropy Loss, optimizer SGD, 0.002 learning rate and 0.9 momentum to train a Fashion-MNIST classifier :
 - 5) [4/45] Change the learning rate :0.0002. Load the 'Prob2.pth and Prob2_1.pth' obtained from Prob 2.3). and Prob 2.4). as the pre-trained model.
 - 6) [4/45] Resume training the model on the Fashion-MNIST dataset.
 - 7) [2/30] Run the testing code and show the accuracy.
 - 8) [8/30] Modify your structure by following setting, train with 6 epochs and test the modified classifier.
 - FC-Layer1: Input: Defined by the third convolutional Layer, Output: 4096
 - FC-Layer2: Input: From FC- Layer1, Output: 2048
 - FC-Layer3: Input: From FC- Layer2, Output: equal to your class size.

Problem 2

2. Prob2.ipynb shows how to train a classifier. Please modify the Prob2.ipynb with the following requirements and set the Cross Entropy Loss, optimizer SGD, 0.002 learning rate and 0.9 momentum to train a Fashion-MNIST classifier :
 - 9) [4/30] By comparing with the two different structures, please describe what you observe according your results.
 - 10) [6/30] Please describe the why we need to use the fully-connected layer and the convolution layer.
 - 11) [3/30] If we change the dataset to CIFAR100 (RGB, 100 classes) from Fashion-MNIST (Gray, 10 classes), what we need to modify in our network.

Solution 2.1

Refer to [Problem 2.A](#)

Design a model with the following structure and complete the table in the next page.

First Conv. layer: Input: Gray scale, Output Channel 32, second Conv. layer:

Output Channel 64, third Conv. layer: Output Channel 128.

FC-Layer1: Input: Defined by the third convolutional Layer, Output: 1000

FC-Layer2: Input: From FC- Layer1, Output: 500

FC-Layer3: Input: From FC- Layer2, Output: equal to your class size

Solution 2.1

Refer to [Problem 2.A](#)

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.relu = nn.ReLU()

        self.conv1 = nn.Conv2d(1, 32, 3, 1, 1)
        self.maxpool1 = nn.MaxPool2d(2, 2, 0)

        self.conv2 = nn.Conv2d(32, 64, 3, 2, 0)
        self.maxpool2 = nn.MaxPool2d(2, 2, 0)

        self.conv3 = nn.Conv2d(64, 128, 4, 2, 1)
        self.avgpool1 = nn.AvgPool2d(2, 1, 1)

        self.fc1 = nn.Linear(128*2*2, 1000)
        self.elu1 = nn.ELU()
        self.fc2 = nn.Linear(1000, 500)
        self.elu2 = nn.ELU()
        self.fc3 = nn.Linear(500, 10)
```

Shape:
(Batch size, 128, 2, 2)

```
def forward(self, x):
    batchsize = x.shape[0]
    x = self.maxpool1(self.relu(self.conv1(x)))
    x = self.maxpool2(self.relu(self.conv2(x)))
    x = self.avgpool1(self.relu(self.conv3(x)))
    # print(x.shape)
    x = x.view(batchsize, -1)
    # print(x.shape)
    x = self.elu1(self.fc1(x))
    x = self.elu2(self.fc2(x))
    x = self.fc3(x)

    return x
```

Solution 2.1

Refer to [Problem 2.A](#)

Shape: (Batch size, 128, 2, 2)

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.relu = nn.ReLU()

        self.conv1 = nn.Conv2d(1, 32, 3, 1, 1)
        self.maxpool1 = nn.MaxPool2d(2, 2, 0)

        self.conv2 = nn.Conv2d(32, 64, 3, 2, 0)
        self.maxpool2 = nn.MaxPool2d(2, 2, 0)

        self.conv3 = nn.Conv2d(64, 128, 4, 2, 1)
        self.avgpool1 = nn.AvgPool2d(2, 1, 1)

        self.fc1 = nn.Linear(128*2*2, 1000)
        self.elu1 = nn.ELU()
        self.fc2 = nn.Linear(1000, 500)
        self.elu2 = nn.ELU()
        self.fc3 = nn.Linear(500, 10)
```

```
def forward(self, x):
    batchsize = x.shape[0]
    x = self.maxpool1(self.relu(self.conv1(x)))
    x = self.maxpool2(self.relu(self.conv2(x)))
    x = self.avgpool1(self.relu(self.conv3(x)))
    # print(x.shape)
    x = x.view(batchsize, -1)
    # print(x.shape)
    x = self.elu1(self.fc1(x))
    x = self.elu2(self.fc2(x))
    x = self.fc3(x)

    return x
```

Net ()

Layer type	Input channel	Output channel	Filter size	Stride	padding
Conv1	A:1	B:32	3	1	1
ReLU					
MaxPool			2	2	0
Conv2	B:32	C:64	3	2	0
ReLU					
MaxPool			2	2	0
Conv3	C:64	D:128	4	2	1
ReLU					
AvgPool			2	1	1
Linear1	E:512	G:1000			
ELU					
Linear2	G:1000	F:500			
ELU					
Linear3	F:500	G:10			

Refer to [Problem 2.A](#)

Please crop the parts that you modify in Prob2.ipynb and paste to the solution.docx.

Solution 2.2

Refer to [Problem 2.E](#)

Train the classifier and show the losses during the training process.

```
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5), (0.5))])

trainset = torchvision.datasets.FashionMNIST(root='./data', train=True, download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4, shuffle=True, num_workers=2)

testset = torchvision.datasets.FashionMNIST(root='./data', train=False, download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4, shuffle=False, num_workers=2)

classes = ('T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat', 'Sandal',
           'Shirt', 'Sneaker', 'Bag', 'Ankle boot')
```

The original codes are
((0.5,0.5,0.5),(0.5,0.5,0.5))

It represents your datasets are 3 channels.
If you want to use the gray-scale dataset,
you need to modify the codes,
((0.5),(0.5))

Solution 2.2

Refer to [Problem 2.E](#)

The original codes are ((0.5,0.5,0.5),(0.5,0.5,0.5))

It represents your datasets are 3 channels. If you want to use the gray-scale dataset, you need to modify the codes to ((0.5),(0.5))

```
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5), (0.5))])

trainset = torchvision.datasets.FashionMNIST(root='./data', train=True, download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4, shuffle=True, num_workers=2)

testset = torchvision.datasets.FashionMNIST(root='./data', train=False, download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4, shuffle=False, num_workers=2)

classes = ('T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat', 'Sandal',
           'Shirt', 'Sneaker', 'Bag', 'Ankle boot')
```

Solution 2.2

```
criterion = nn.CrossEntropyLoss()  
optimizer = optim.SGD(net.parameters(), lr=0.002, momentum=0.9)
```

```
for epoch in range(3): # loop over the dataset multiple times  
  
    running_loss = 0.0  
    for i, data in enumerate(trainloader):  
        # get the inputs; data is a list of [inputs, labels]  
        inputs, labels = data  
        inputs = inputs.cuda()  
        labels = labels.cuda()  
  
        # zero the parameter gradients  
        optimizer.zero_grad()  
  
        # forward + backward + optimize  
        outputs = net(inputs)  
        loss = criterion(outputs, labels)  
        loss.backward()  
        optimizer.step()
```

Refer to
[Exercise 2-4 \(P.61\)](#)
[Exercise 2-4 \(P.62\)](#)

```
Epoch : 1 steps : 1000 Training Loss : 2.127310976743698  
Epoch : 1 steps : 2000 Training Loss : 1.0009411244057118  
Epoch : 1 steps : 3000 Training Loss : 0.7320705399112776  
Epoch : 1 steps : 4000 Training Loss : 0.6135896655730904  
Epoch : 1 steps : 5000 Training Loss : 0.570787304927595
```

```
Epoch : 3 steps : 10000 Training Loss : 0.3041461807802189  
Epoch : 3 steps : 11000 Training Loss : 0.29006112643110193  
Epoch : 3 steps : 12000 Training Loss : 0.27181791267670635  
Epoch : 3 steps : 13000 Training Loss : 0.2741309449586397  
Epoch : 3 steps : 14000 Training Loss : 0.2826621224831997  
Epoch : 3 steps : 15000 Training Loss : 0.2622677051232313  
Finished Training
```

1. Feed the input into the model and get the prediction.
2. Use the defined loss function to calculate the loss between the prediction and the label.
3. Use `backward()` to compute the gradient and use the `optimizer.step()` to update the weight

Solution 2.2

Refer to
Exercise 2-4 (P.64)

Testing

```
▶ correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images = images.cuda()
        labels = labels.cuda()
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy : %d %%' % (100 * correct / total))

☞ Accuracy : 89 %
```

1. Use for loop to get the entire testing data
2. Feed the testing data and calculate the accuracy

Solution 2.3 & 2.4

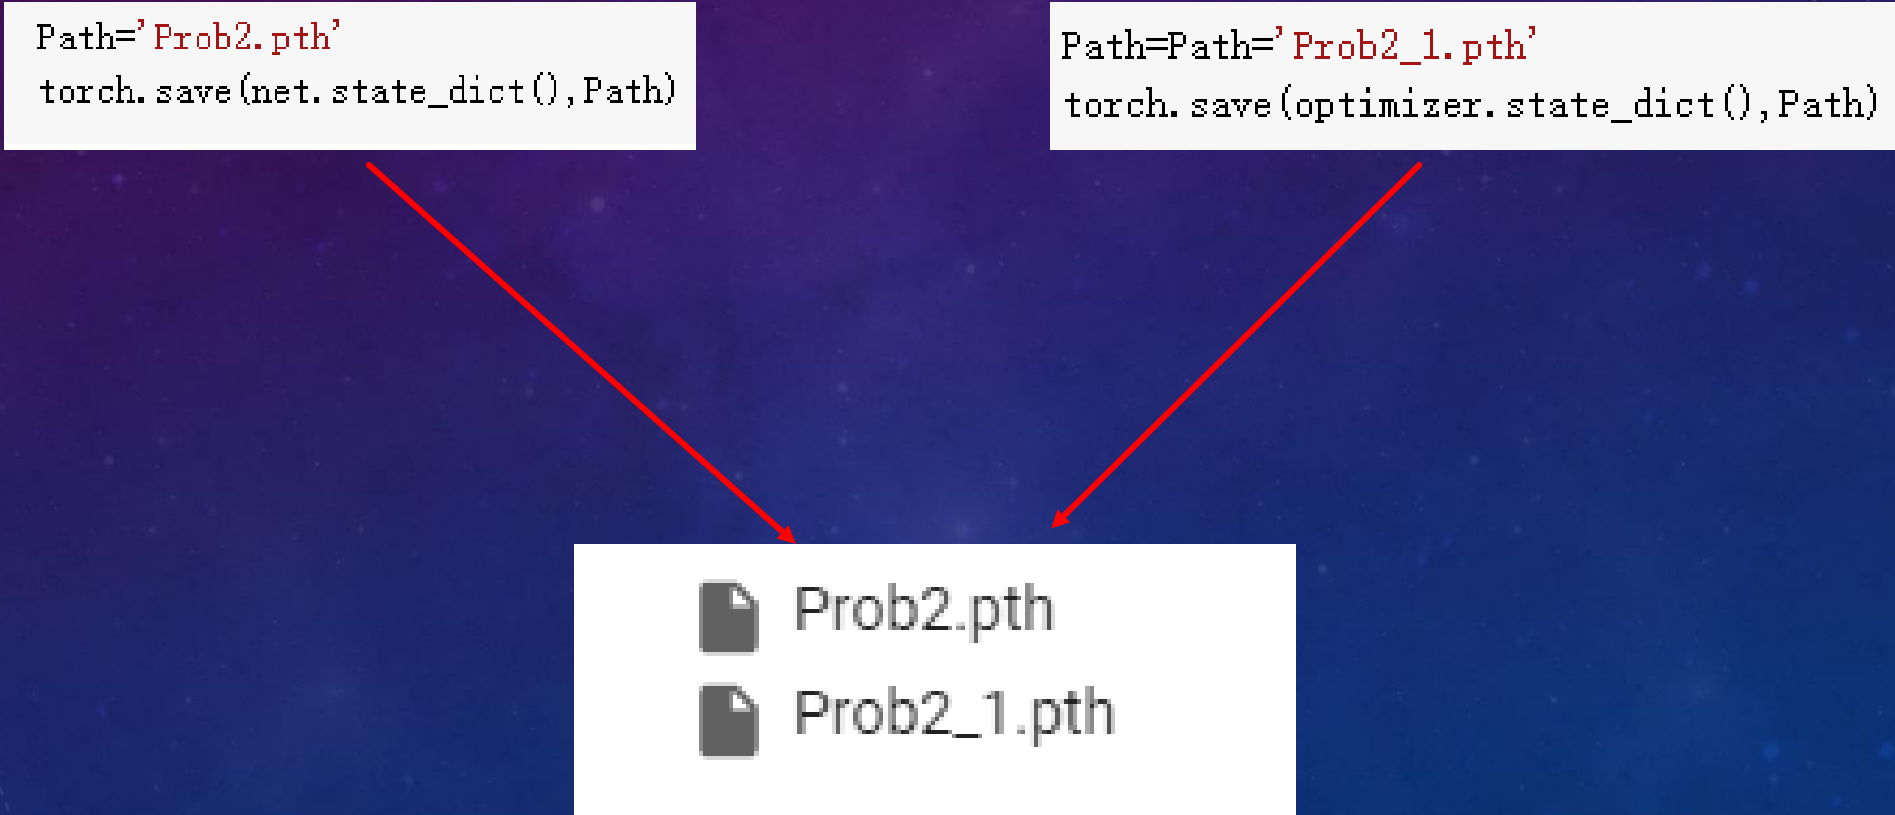
Refer to
[Problem 2.B & 2.C](#)

Save the model and name it as 'Prob2.pth'.

Save the optimizer and name it as 'Prob2_1.pth'.

```
Path='Prob2.pth'  
torch.save(net.state_dict(),Path)
```

```
Path=Path='Prob2_1.pth'  
torch.save(optimizer.state_dict(),Path)
```



Prob2.pth
Prob2_1.pth

Solution 2.5

Refer to [Problem 2.F](#)

Change the learning rate :0.0002. Load the 'Prob2.pth and Prob2_1.pth' obtained from Prob 2.3). and Prob 2.4). as the pre-trained model.

```
net = Net()
checkpoint = 'Prob2.pth'
checkpoint = torch.load(checkpoint)
net.load_state_dict(checkpoint)
net = net.cuda()

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.0002, momentum=0.9)

checkpoint = 'Prob2_1.pth'
checkpoint = torch.load(checkpoint)
optimizer.load_state_dict(checkpoint)
```

Solution 2.6

Refer to
Exercise 2-4 (P.62)

Resume training the model on the Fashion-MNIST dataset.

```
for epoch in range(3):    # loop over the dataset multiple times

    running_loss = 0.0
    for i, data in enumerate(trainloader):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data

        inputs = inputs.cuda()
        labels = labels.cuda()

        # zero the parameter gradients
        optimizer.zero_grad()
        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
```

```
Epoch : 1 steps : 1000 Training Loss : 0.2773041429291975
Epoch : 1 steps : 2000 Training Loss : 0.2512894716824194
Epoch : 1 steps : 3000 Training Loss : 0.2637920055533687
Epoch : 1 steps : 4000 Training Loss : 0.26050553433461393
Epoch : 1 steps : 5000 Training Loss : 0.2440475834285321
Epoch : 1 steps : 6000 Training Loss : 0.25720433985322777
Epoch : 1 steps : 7000 Training Loss : 0.25939772264012206
```

```
Epoch : 3 steps : 7000 Training Loss : 0.2087033114988908
Epoch : 3 steps : 8000 Training Loss : 0.2131072640818001
Epoch : 3 steps : 9000 Training Loss : 0.20193223340443522
Epoch : 3 steps : 10000 Training Loss : 0.23447983676543846
Epoch : 3 steps : 11000 Training Loss : 0.21496245267047742
Epoch : 3 steps : 12000 Training Loss : 0.23314484732411528
```

Solution 2.7

Refer to
[Exercise 2-4 \(P.64\)](#)

Testing

```
|  
[ ] dataiter = iter(testloader)  
    images, labels = dataiter.next()  
  
[ ] correct = 0  
    total = 0  
    with torch.no_grad():  
        for data in testloader:  
            images, labels = data  
            images, labels = images.cuda(), labels.cuda()  
            outputs = net(images)  
            _, predicted = torch.max(outputs.data, 1)  
            total += labels.size(0)  
            correct += (predicted == labels).sum().item()  
  
    print('Accuracy : %d %%' % (100 * correct / total))  
  
Accuracy : 90 %
```

Solution 2.8

Refer to [Problem 2.A](#)

Modify your structure by following setting, train with 6 epochs and test the modified classifier.

- FC-Layer1: Input: Defined by the third convolutional Layer, Output: 4096
- FC-Layer2: Input: From FC- Layer1, Output: 2048
- FC-Layer3: Input: From FC- Layer2, Output: equal to your class size.

```
class Net2(nn.Module):
    def __init__(self):
        super(Net2, self).__init__()
        self.relu = nn.ReLU()

        self.conv1 = nn.Conv2d(1, 32, 3, 1, 1)
        self.maxpool1 = nn.MaxPool2d(2, 2, 0)

        self.conv2 = nn.Conv2d(32, 64, 3, 2, 0)
        self.maxpool2 = nn.MaxPool2d(2, 2, 0)

        self.conv3 = nn.Conv2d(64, 128, 4, 2, 1)
        self.avgpool1 = nn.AvgPool2d(2, 1, 1)

        self.fc1 = nn.Linear(128*2*2, 4096)
        self.elu1 = nn.ELU()
        self.fc2 = nn.Linear(4096, 2048)
        self.elu2 = nn.ELU()
        self.fc3 = nn.Linear(2048, 10)
```


Solution 2.8

Refer to
Exercise 2-4 (P.62)

```
net2 = Net2()
net2 = net2.cuda()
```

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net2.parameters(), lr=0.002, momentum=0.9)
```

```
for epoch in range(6): # loop over the dataset multiple times

    running_loss = 0.0
    for i, data in enumerate(trainloader):
        # get the inputs: data is a list of [inputs, labels]
        inputs, labels = data
        inputs = inputs.cuda()
        labels = labels.cuda()

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net2(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
```

Solution 2.8

Refer to
Exercise 2-4 (P.62)

```
net2 = Net2()  
net2 = net2.cuda()
```

Initialize the net2 parameter and transfer it to GPU mode

```
criterion = nn.CrossEntropyLoss()  
optimizer = optim.SGD(net2.parameters(), lr=0.002, momentum=0.9)
```

Use SGD optimizer to update the parameters in the Net2.

Note:

Remember to use the net2.parameter instead of using net.parameter.

Solution 2.8

Refer to
Exercise 2-4 (P.62)

```
for epoch in range(6):    # loop over the dataset multiple times

    running_loss = 0.0
    for i, data in enumerate(trainloader):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
        inputs = inputs.cuda()
        labels = labels.cuda()

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net2(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
```

Remember to modify
the codes for feeding
data into the net2.

Solution 2.9 & 2.10

- By comparing with the two different structures, please describe what you observe according your results
- Solution 2.9:
 - More iteration times is needed to achieve the similar performance when enlarging the FC layer. It means that the modification will increase the complexities. However, the performance does not approve the better results in this modification. The fact shows that it is a batter way to design the appropriate FC layers according to the dataset.
- Please describe the why we need to use the fully-connected layer and the convolution layer.
- Solution 2.10:
 - CNNs have two main parts: A convolution mechanism that breaks up the images into features and analyzes them. A fully connected layer that takes the outputs of convolution and predicts the best answer to describe the image.
 - Convolution is used in deep neural networks, which can reduce parameters through parameter sharing and sparse connections. When using fully connected layer to connect two large dimension features, there will be a lot of parameters.

Solution 2.11

If we change the dataset to CIFAR100 (RGB, 100 classes) from Fashion-MNIST (Gray, 10 classes), what we need to modify in our network.

CIFAR100: RGB, 100 classes

Fashion-MNIST: Gray, 10 classes

When changing the dataset, we need to modify the first convolution layer, the first and third fully connected layers setting to match the dataset format. (i.e., image size, channels, classes)

Problem 3 [15/100]

Refer to [Problem 3](#)

3. The output dimensions of Conv4 is $128 \times 96 \times 56$ (i.e., Channel \times Width \times Height) , please calculate the dimensions of the inputs and the outputs from Conv1, Conv2, Conv3, Conv5.

Layer type	Input channel	Output channel	Filter size	Stride	Padding
Conv1	3	16	3	1	1
AvgPool			4	1	0
Conv2	16	16	4	2	0
MaxPool			2	2	0
Conv3	16	32	2	1	1
MaxPool			2	1	0
Conv4	32	128	3	2	2
AvgPool			3	1	1
Conv5	128	256	7	1	0

Solution3

$$\text{Output} = \frac{\text{Input} - \text{kernel size} + 2 \times \text{Paddiing}}{\text{Stride}} + 1$$

$$\text{Input} = (\text{Output} - 1) \times \text{Stride} - 2 \times \text{Paddiing} + \text{kernel size}$$

- Input-> [3,761,441]
- Conv1-> [8, 761, 441]
- AvgPool1-> [8, 758, 438]
- Conv2-> [16, 378, 218]
- MaxPool1-> [16, 189, 109]
- Conv3-> [32, 190, 110]
- MaxPool2-> [32, 189, 109]
- Conv4-> [128, 96, 56]

$$761 = (761 - 1) \times 1 - 2 \times 1 + 3, \quad 441 = (2065 - 1) \times 1 - 2 \times 1 + 3$$

$$761 = (758 - 1) \times 1 - 2 \times 0 + 4, \quad 441 = (438 - 1) \times 1 - 2 \times 0 + 4$$

$$758 = (378 - 1) \times 2 - 2 \times 0 + 4, \quad 438 = (218 - 1) \times 2 - 2 \times 0 + 4$$

$$378 = (189 - 1) \times 2 - 2 \times 0 + 2, \quad 218 = (515 - 1) \times 2 - 2 \times 0 + 2$$

$$189 = (190 - 1) \times 1 - 2 \times 1 + 2, \quad 109 = (110 - 1) \times 1 - 2 \times 1 + 2$$

$$190 = (189 - 1) \times 1 - 2 \times 0 + 2, \quad 110 = (109 - 1) \times 1 - 2 \times 0 + 2$$

$$189 = (96 - 1) \times 2 - 2 \times 2 + 3, \quad 109 = (56 - 1) \times 2 - 2 \times 2 + 3$$

Refer to [Problem 3](#)

Solution3

- AvgPool2->[128, 96, 56]
- Conv5-> [256, 90, 50]

$$96 = \frac{96 - 3 + 2 \times 1}{1} + 1,$$

$$90 = \frac{96 - 7 + 2 \times 0}{1} + 1,$$

Refer to Problem 3

$$56 = \frac{56 - 3 + 2 \times 1}{1} + 1$$

$$50 = \frac{56 - 7 + 2 \times 1}{1} + 1$$

Reference


Content Overview

- The examples and the exercises
- The solutions to the sample problems

The Examples and The Exercises

Example 2-1: Feature Map Visualization

- Please download the “2-1_Feature_map_visualization.zip” from the Moodle, which is built on the VGG-16 trained on the ImageNet.
- Upload the 2-1_Feature_map_visualization.ipynb and imagenet1000_clsidx_to_labels.txt to the Google Colab.
- Upload the “jellyfish.jpg” to the Google Colab.
- Run the codes and get the feature map from the selected layer.



Example 2-1: Feature Map Visualization

imagenet1000_clsidx_to_labels.txt


(in “Feature_map_visualization_v2.7z”)

```
0: 'tench, Tinca tinca',
1: 'goldfish, Carassius auratus',
2: 'great white shark, white shark, man-eater, man-eating shark, Carcharodon carcharias',
3: 'tiger shark, Galeocerdo cuvieri',
4: 'hammerhead, hammerhead shark',
5: 'electric ray, crampfish, numbfish, torpedo',
6: 'stingray',
7: 'cock',
8: 'hen',
9: 'ostrich, Struthio camelus',
10: 'brambling, Fringilla montifringilla',
11: 'goldfinch, Carduelis carduelis',
12: 'house finch, linnet, Carpodacus mexicanus',
13: 'junco, snowbird',
14: 'indigo bunting, indigo finch, indigo bird, Passerina cyanea',
15: 'robin, American robin, Turdus migratorius',
16: 'bulbul',
17: 'jay',
18: 'magpie',
19: 'chickadee',
20: 'water ouzel, dipper',
21: 'kite',
22: 'bald eagle, American eagle, Haliaeetus leucocephalus',
23: 'vulture',
24: 'great grey owl, great gray owl, Strix nebulosa',
25: 'European fire salamander, Salamandra salamandra',
26: 'common newt, Triturus vulgaris',
27: 'eft',
28: 'spotted salamander, Ambystoma maculatum',
29: 'axolotl, mud puppy, Ambystoma mexicanum',
30: 'bullfrog, Rana catesbeiana',
31: 'tree frog, tree-frog',
32: 'tailed frog, bell toad, ribbed toad, tailed toad, Ascaphus trui',
33: 'loggerhead, loggerhead turtle, Caretta caretta',
```

Example 2-1 : Feature Map Visualization



Original image: jellyfish.jpg



Probability of the classes

Probability TOP-3:

```
TOP_1
Probability:0.9990069270133972
Predicted: 'jellyfish'

TOP_2
Probability:0.0008054533391259611
Predicted: 'isopod'


TOP_3
Probability:8.906585571821779e-05
Predicted: 'chambered nautilus'
```

Predicted class : jellyfish

Example 2-1 : Feature Map Visualization



Original image: jellyfish.jpg



Example 2-1 : Coding Explanation

Overview of the sample code:

- Main Process for executing
- Function - FeatureVisualization
 - “__init__ (i.e. initialization)” for setting the pretrained model i.e., vgg16 on ImageNet
 - “process_image” for the image preprocessing.
 - “get_multi_feature” for getting the feature maps.
 - “save_feature_to_img” for saving the feature maps.
 - “Predict” for getting the prediction from the given image.

Example 2-1 : Coding Explanation

[Return solution 1.1](#)

Main Process for executing

```
if __name__=='__main__':  
    # get class  
    c = {}  
    with open("imagenet1000_clsid_to_labels.txt") as f:  
        for line in f:  
            (key, val) = line.split(":")  
            c[int(key)] = val.split(",")[0]  
    # Define image path and select the layer  
    myClass=FeatureVisualization('./jellyfish.jpg', 5)  
    print(myClass.pretrained_model2)  
    myClass.save_feature_to_img()  
    myClass.predict()
```

Open the txt file to get the information
(Hint: You need to upload the file to Colab., or you will get the error "No such file.....")

Select the layer 5

Print the Network Architecture

1. Call the function to save the extracted feature from the selected layer in the VGG16.
2. Call the predict function to get the prediction from the given image

Example 2-1 : Coding Explanation

Feature Visualization

```
def __init__(self, img_path, selected_layer):  
    self.img_path=img_path  
    self.selected_layer=selected_layer  
    # Load pretrained model  
  
    self.pretrained_model = models.vgg16(pretrained=True).features  
    self.pretrained_model.eval()  
  
    self.pretrained_model2 = models.vgg16(pretrained=True)  
    self.pretrained_model2.eval()
```

Call the **feature part** of vgg16 pretrained model.
“eval()” is for fixing the pretrained weight.

Call the **entire** vgg16 pretrained model (i.e. **the feature part and classifier part**)
“eval()” is for fixing the pretrained weight.



FC8

(1000)

1000
Classes
Result

Feature Extraction

Example 2-1 : Coding Explanation

```
def process_image(self):  
    img=cv2.imread(self.img_path)  
    img=preprocess_image(img)  
    return img
```

Read the given image and preprocess it before feeding it into the model.

```
def preprocess_image(cv2im, resize_im=True):
```

```
    # Resize image  
    1. if resize_im:  
        cv2im = cv2.resize(cv2im, (224, 224))  
    im_as_arr = np.float32(cv2im)  
    2. im_as_arr = np.ascontiguousarray(im_as_arr[..., ::-1])  
    im_as_arr = im_as_arr.transpose(2, 0, 1) # Convert array to D,W,H  
    # Normalize the channels  
    3. for channel, _ in enumerate(im_as_arr):  
        im_as_arr[channel] /= 255  
    # Convert to float tensor  
    im_as_ten = torch.from_numpy(im_as_arr).float()  
    # Add one more channel to the beginning. Tensor shape = 1,3,224,224  
    4. im_as_ten.unsqueeze_(0)  
    # Convert to Pytorch variable  
    im_as_var = Variable(im_as_ten, requires_grad=True)  
    return im_as_var
```

Preprocess:

1. Resize to the 224x224 (i.e. **VGG16 input size**)
2. Convert the dimension to match the format of PyTorch.
3. Normalize the value of the data (**From 0 to 1, i.e., divide data by 255**)
4. Convert the data type to PyTorch tensor type

Example 2-1 : Coding Explanation

[Return solution 1.2](#)

```
def get_feature(self):  
    # Image preprocessing  
    input=self.process_image()  
    #print("input.shape: {}".format(input.shape))  
    x=input  
    for index,layer in enumerate(self.pretrained_model):  
        x=layer(x)  
        #print("x: {}".format(x.shape))  
        if (index == self.selected_layer):  
            return x
```

Feed the preprocessed image into the feature part of VGG16 model to extract the feature from the given layer.
(i.e., the index of layer equal to the given value)

```
def get_multi_feature(self):  
    # Get the feature map  
    features=self.get_feature()  
    #print(features.shape)  
    result_path = './feat_first' + str(self.selected_layer)  
  
    if not os.path.exists(result_path):  
        os.makedirs(result_path)  
    print("On layer:{}, We can get the {} feature maps".format(self.selected_layer, features.shape[1]))  
    #print(features.shape[1])  
    for i in range(features.shape[1]):  
        feature=features[:,i,:]  
        feature=feature.view(feature.shape[1],feature.shape[2])  
        feature = feature.data.numpy()  
        feature = 1.0 / (1 + np.exp(-1 * feature))  
        feature = np.round(feature * 255)  
        save_name = result_path + '/' + str(i) + '.jpg'  
        cv2.imwrite(save_name, feature)
```

1. Create the folder to save the feature map.
2. Use for loop to save the extracted feature maps

Example 2-1 : Coding Explanation

```
def save_feature_to_img(self):  
    #to numpy  
    feature=self.get_single_feature()  
    self.get_multi_feature()  
    feature=feature.data.numpy()  
  
    #use sigmod to [0,1]  
    # print(feature[0])  
    feature= 1.0/(1+np.exp(-1*feature))  
  
    # to [0,255]  
    feature=np.round(feature*255)  
    #print(self.selected_layer)  
    save_name = './feat_first' + str(self.selected_layer) + '.jpg'  
    cv2.imwrite(save_name, feature)
```

Call the function of extracting feature maps

Save the sample feature map

Example 2-1 : Coding Explanation

```
def predict(self):  
    input=self.process_image()  
    outputs = self.pretrained_model2(input)  
  
    s = torch.nn.Softmax(dim=1)  
    result = s(outputs)  
    self.plot_probablity(result)  
  
    prob, predicted = result.sort(1,descending=True)  
    prob = prob.data.numpy()  
  
    predicted = predicted.data.numpy()  
  
    print("Probablity TOP-3:\n")  
    print("")  
    for i in range(3):  
        print("TOP_" +str(i+1))  
        print("Probablity: {}".format(prob[0][i]))  
        print("Predicted: {}\n".format(c[int(predicted[0][i])]))  
    return outputs
```

Call the preprocessed data, feed it into the entire vgg16 pretrained model, and get the output from the classifier.

Call the **Softmax** function to transform the output value to the **probability** and plot the figure.

Sort the predicted probability and show the first three value and its class in the ImageNet .

Example 2-2 : Feature Map Visualization

- Please download the “2-2_Feature_map_visualization.zip” on the Moodle and choose your own images from Internet.
- Upload the 2-2_Feature_map_visualization.ipynb and imagenet1000_clsidx_to_labels.txt to the Google Colab.
- Upload the given images “g1.jpg” and “g2.jpg” to the Google Colab.
- Run the codes and get the probabilities of these images.



Example 2-3: Feature Comparison

- Please download the “2-2_Feature_map_visualization.zip” on the Moodle
- Upload the 2-2_Feature_map_visualization.ipynb and imagenet1000_clsidx_to_labels.txt to the Google Colab.
- Upload the given images “g3.jpg” and “g4.jpg” to the Google Colab.
- Run the codes and get the comparisons.



Example 2-2 & 2-3: Feature Map Visualization

```
if __name__ == '__main__':
    # get class
    c = {}
    with open("imagenet1000_clsidx_to_labels.txt") as f:
        for line in f:
            (key, val) = line.split(":")
            c[int(key)] = val.split(",")[0]

    # Define image path and select the layer
    myClass=FeatureVisualization('./dog6.jpg',12)
    Compare=FeatureVisualization('./dog9.jpg',12)
    print(myClass.pretrained_model2)

    myClass.save_feature_to_img()
    Compare.save_feature_to_img1()
    print("The first picture classification predict:")
    myClass_vector = myClass.predict()
    print("The second picture classification predict:")
    Compare_vector = Compare.predict()
    #Define cosine similarity
    cos= nn.CosineSimilarity(dim=1)
    #Define Euclidean distance
    euclidean_dist = torch.dist(myClass_vector, Compare_vector, p=2)
    cosine_dist = 1-cos(myClass_vector, Compare_vector)
    print(' Verification:')
    if cosine_dist < 0.6:
        print("They are the same!")
        print("Their cosine_distance: {}".format(cosine_dist))
    else:
        print("They are not the same!")
        print("Their cosine_distance: {}".format(cosine_dist))

    print("Their euclidean_dist: {}".format(euclidean_dist))
```

Return solution 1.2

solution 1.4, solution 1.5

solution 1.6, solution 1.7

solution 1.8

← Define the Cosine Similarity function

← Calculate the Euclidean distance between different pictures
Calculate the Cosine distance between different pictures

← Define the threshold

Example 2-2 & 2-3 : Feature Overview

Results:

On layer:2, We can get the 64 feature maps
The first picture classification predict:
Probability TOP-3:

TOP_1
Probability:0.9882549047470093
Predicted: 'jellyfish'

TOP_2
Probability:0.00702690239995718
Predicted: 'isopod'

TOP_3
Probability:0.0019321587169542909
Predicted: 'nematode'



The first picture

The second picture classification predict:
Probability TOP-3:

TOP_1
Probability:0.2607852518558502
Predicted: 'sports car'

TOP_2
Probability:0.20074793696403503
Predicted: 'beach wagon'

TOP_3
Probability:0.13690434396266937
Predicted: 'convertible'

Verification:
They are not the same!
Their cosine_similarity:tensor([0.0470], grad_fn=<DivBackward0>)
Their euclidean_dist:135.32333374023438



The second picture

Exercise 2-2: Feature Map Visualization

- Please download the “2-2_Feature_map_visualization.zip” on the Moodle and choose your own images from Internet.
- Upload the 2-2_Feature_map_visualization.ipynb and imagenet1000_clsidx_to_labels.txt to the Google Colab.
- Compare the probability of the images that contain multi classes and different variations (pose, occlusion, age).
- Please write down results and your codes in MS Word to the Moodle



Exercise 2-2 & 2-3: Feature Map Visualization

[Return solution 1.3](#)

```
if __name__ == '__main__':
    # get class
    c = {}
    with open("imagenet1000_clsidx_to_labels.txt") as f:
        for line in f:
            (key, val) = line.split(":")
            c[int(key)] = val.split(",")[0]

    # Define image path and select the layer
    myClass=FeatureVisualization('./dog6.jpg',12)
    Compare=FeatureVisualization('./dog9.jpg',12)
    print(myClass.pretrained_model2)

    myClass.save_feature_to_img()
    Compare.save_feature_to_img1()

    print("The first picture classification predict:")
    myClass_vector = myClass.predict()
    print("The second picture classification predict:")
    Compare_vector = Compare.predict()

    #Define cosine similarity
    cos= nn.CosineSimilarity(dim=1)
    #Define Euclidean distance
    euclidean_dist = torch.dist(myClass_vector, Compare_vector, p=2)
    cosine_dist = 1-cos(myClass_vector, Compare_vector)
    print("Verification:")
    if cosine_dist < 0.6:
        print("They are the same!")
        print("Their cosine_distance: {}".format(cosine_dist))
    else:
        print("They are not the same!")
        print("Their cosine_distance: {}".format(cosine_dist))

    print("Their euclidean_dist: {}".format(euclidean_dist))
```

← Show the predicted classes and probabilities

Exercise 2-2 & 2-3: Feature Map Visualization



TOP_1
Probability:0.730004072189331
Predicted: 'jellyfish'

TOP_2
Probability:0.055244747549295425
Predicted: 'cup'

TOP_3
Probability:0.023521317169070244
Predicted: 'vase'



TOP_1
Probability:0.773815393447876
Predicted: 'Samoyed'

TOP_2
Probability:0.09530658274888992
Predicted: 'West Highland white terrier'

TOP_3
Probability:0.014327057637274265
Predicted: 'komondor'

Exercise 2-2 & 2-3: Feature Map Visualization



```
They are not the same!  
Their cosine_distance:tensor([0.9956], grad_fn=<RsubBackward1>)  
Their euclidean_dist:99.01476287841797
```

The cosine similarity is : $1 - 0.9956 = 0.0044$

[solution 1.4](#), [solution 1.5](#)
[solution 1.6](#), [solution 1.7](#)
[solution 1.8](#)

Exercise 2-4 – Build A Classifier

- Please download the “2-4_CIFAE10.ipynb” on the Moodle, run the sample code, and change the following parameters:
 - Epoch
 - Learning Rate: 0.1, 0.01, 0.001
- Please upload your result and observations in MS Words to the Moodle.

Exercise 2-4 – Build A Classifier

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

Define a Convolutional Neural Network

Exercise 2-4 – Build A Classifier

Preprocess function

```
transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4, shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4, shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

```
net = Net()
```

Initialize the network

Define the dataset and its classes

Exercise 2-4 – Build A Classifier

[Return to
solution 2.2](#)

```
criterion = nn.CrossEntropyLoss()  
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

Change the learning rate: 0.1, 0.01 and 0.001

Note: If you want to use Adam optimizer,
betas = {0.9 and 0.999}, the codes are:

```
criterion = nn.CrossEntropyLoss()  
optimizer = optim.Adam(net.parameters(), lr=0.001, betas=(0.9, 0.999))
```

1. Define the Cross Entropy Loss
2. Define the SGD optimizer which the learning rate is 0.001 and the momentum is 0.9.

1. Define the Cross Entropy Loss
2. Define the Adam optimizer which the learning rate is 0.001 and the betas are 0.9 and 0.999.

Exercise 2-4 – Build A Classifier

[Return to
solution 2.2, 2.8](#)

Epoch number

```
for epoch in range(3): # loop over the dataset multiple times

    running_loss = 0.0
    for i, data in enumerate(trainloader, 0):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
        if i % 2000 == 1999: # print every 2000 mini-batches

            print("Epoch : {} steps : {} Training Loss : {}".format(epoch + 1, i + 1, running_loss / 2000) )
            running_loss = 0.0
            save_checkpoint({'net':net.state_dict()}, 'test_epoch{}'.format(epoch+1))

    print('Finished Training')
```

1. Feed the input into the model and get the prediction.
2. Use the defined loss function to calculate the loss between the prediction and the label.
3. Use backward() to compute the gradient and use the optimizer.step() to update the weight

Exercise 2-4 – Build A Classifier

The training process of different learning rate:

```
Epoch : 1 steps : 2000 Training Loss : 2.357952641606331
Epoch : 1 steps : 4000 Training Loss : 2.358549834549427
Epoch : 1 steps : 6000 Training Loss : 2.363395507276058
Epoch : 1 steps : 8000 Training Loss : 2.3620128685832023
Epoch : 1 steps : 10000 Training Loss : 2.3548867295980456
Epoch : 1 steps : 12000 Training Loss : 2.3602904160022735
Epoch : 2 steps : 2000 Training Loss : 2.364006470501423
Epoch : 2 steps : 4000 Training Loss : 2.3610252693295477
Epoch : 2 steps : 6000 Training Loss : 2.3646557998657225
Epoch : 2 steps : 8000 Training Loss : 2.3612379420399665
Epoch : 2 steps : 10000 Training Loss : 2.359113136589527
Epoch : 2 steps : 12000 Training Loss : 2.361136519730091
Epoch : 3 steps : 2000 Training Loss : 2.360108473300934
Epoch : 3 steps : 4000 Training Loss : 2.3645326865315437
Epoch : 3 steps : 6000 Training Loss : 2.3571521565318108
Epoch : 3 steps : 8000 Training Loss : 2.3616783508062364
Epoch : 3 steps : 10000 Training Loss : 2.356024751186371
Epoch : 3 steps : 12000 Training Loss : 2.360763477861881
Finished Training
```

Learning rate: 0.1

```
Epoch : 1 steps : 2000 Training Loss : 2.0899574621915815
Epoch : 1 steps : 4000 Training Loss : 1.9607795716822147
Epoch : 1 steps : 6000 Training Loss : 1.9563240223526954
Epoch : 1 steps : 8000 Training Loss : 1.957445238739252
Epoch : 1 steps : 10000 Training Loss : 1.9918364935815334
Epoch : 1 steps : 12000 Training Loss : 1.9577875487208367
Epoch : 2 steps : 2000 Training Loss : 2.011085530459881
Epoch : 2 steps : 4000 Training Loss : 2.0082346482574938
Epoch : 2 steps : 6000 Training Loss : 2.0100900876820087
Epoch : 2 steps : 8000 Training Loss : 2.0044543738663196
Epoch : 2 steps : 10000 Training Loss : 1.969518426090479
Epoch : 2 steps : 12000 Training Loss : 1.9845602488517762
Epoch : 3 steps : 2000 Training Loss : 1.9992908894717694
Epoch : 3 steps : 4000 Training Loss : 2.0016448673307896
Epoch : 3 steps : 6000 Training Loss : 2.016233505010605
Epoch : 3 steps : 8000 Training Loss : 2.028977326095104
Epoch : 3 steps : 10000 Training Loss : 2.01645450925827
Epoch : 3 steps : 12000 Training Loss : 2.056691348493099
Finished Training
```

Learning rate: 0.01

```
Epoch : 1 steps : 2000 Training Loss : 2.1889700249433517
Epoch : 1 steps : 4000 Training Loss : 1.8390005451440812
Epoch : 1 steps : 6000 Training Loss : 1.643870963960886
Epoch : 1 steps : 8000 Training Loss : 1.5679095338881015
Epoch : 1 steps : 10000 Training Loss : 1.492728895097971
Epoch : 1 steps : 12000 Training Loss : 1.4820199556872249
Epoch : 2 steps : 2000 Training Loss : 1.407832405924797
Epoch : 2 steps : 4000 Training Loss : 1.3714569466710091
Epoch : 2 steps : 6000 Training Loss : 1.3651220782622695
Epoch : 2 steps : 8000 Training Loss : 1.3291044723726808
Epoch : 2 steps : 10000 Training Loss : 1.284140573028475
Epoch : 2 steps : 12000 Training Loss : 1.2962907982245087
Epoch : 3 steps : 2000 Training Loss : 1.2023748714327813
Epoch : 3 steps : 4000 Training Loss : 1.2107961179297417
Epoch : 3 steps : 6000 Training Loss : 1.1705148331448436
Epoch : 3 steps : 8000 Training Loss : 1.2100771391429006
Epoch : 3 steps : 10000 Training Loss : 1.185110432397574
Epoch : 3 steps : 12000 Training Loss : 1.1803923727944494
Finished Training
```

Learning rate: 0.001

Exercise 2-4 – Build A Classifier

[Return to
solution 2.2](#)

```
dataiter = iter(testloader)
images, labels = dataiter.next()
```

Use next() to get 1 batch data as a test sample.

```
outputs = net(images)
_, predicted = torch.max(outputs, 1)
```

```
print('Predicted: ', ' '.join('%5s' % classes[predicted[j]] for j in range(4)))
correct = 0
total = 0
```

```
with torch.no_grad():
    for data in testloader:
        images, labels = data
        outputs = net(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print('Accuracy : %d %%' % (100 * correct / total))
```

1. Feed the input data into the model and get the value of CIFAR10's 10 classes.
2. Use torch.max to get the maximum score and its index

1. Use for loop to get the entire testing data
2. Feed the testing data and calculate the accuracy

Accuracy : 10 %

Learning rate: 0.1

Accuracy : 25 %

Learning rate: 0.01

Accuracy : 61 %

Learning rate: 0.001

The background is a deep blue gradient with a subtle pattern of white stars and nebulae. Overlaid on this are several faint, white technical diagrams. In the top right, there is a large circular scale with degree markings from 90 to 210 and concentric circles with arrows indicating rotation. In the bottom right, there is a diagram with concentric circles and dashed lines, possibly representing a signal or data flow. In the bottom left, there is a partial view of a circular diagram with an arrow. The title text is centered in the middle of the slide.

Solutions To The Sample Problems

Problem 1 [30/100]

1. Prob1.ipynb gives you a VGG-16 trained on ImageNet. Upload the “imagenet1000_clsidx_to_labels.txt “ and g1.jpg and g2.jpg to the Colab. Use Prob1.ipynb to show the following:
 - A. The feature maps and dimensions extracted from Layer 10. [8/30] (Example 2-1, Page 10)
 - B. Calculate the Euclidean distance between the images g1.jpg and g2.jpg, which are given with the code. [6/30] (Example 2-2 & 2-3, Page 19)
 - C. Please list the changes of dimension when feeding a image to the VGG-16. [16/30] (Example 2-1, Page 10)

Solution 1A

Return solution 1.1

```
if __name__=='__main__':  
    # get class  
    c = {}  
    with open("imagenet1000_clsidx_to_labels.txt") as f:  
        for line in f:  
            (key, val) = line.split(":")  
            c[int(key)] = val.split(",")[0]  
    # Define image path and select the layer  
    first=FeatureVisualization('./g1.jpg',10)  
    second=FeatureVisualization('./g2.jpg',10)
```

The feature maps and dimensions extracted from Layer 10.

```
def get_feature(self):  
    # Image preprocessing  
    input=self.process_image()  
    #print("input.shape: {}".format(input.shape))  
    x=input  
    for index,layer in enumerate(self.pretrained_model):  
        x=layer(x)  
        print("x: {}".format(x.shape))  
        if (index == self.selected_layer):  
            return x
```

x:torch.Size([1, 256, 56, 56])
On layer:10, We can get the 256 feature maps

Solution 1B

Return solution 1.2

```
first.save_feature_to_img()
second.save_feature_to_img1()
print("The first picture classification predict:")
first_vector = first.predict()
print("The second picture classification predict:")
second_vector = second.predict()

#Define Euclidean distance
euclidean_dist = torch.dist(first_vector, second_vector, p=2)

print("Verification:")
print("Their euclidean_dist: {}".format(euclidean_dist))
```

Define the Euclidean distance to calculate the distance between the given images.

```
Verification:
Their euclidean_dist:84.86775970458984
```

Solution 1C

```
def get_feature(self):  
    # Image preprocessing  
    input=self.process_image()  
    #print("input.shape: {}".format(input.shape))  
    x=input  
    for index,layer in enumerate(self.pretrained_model):  
        x=layer(x)  
        print("x: {}".format(x.shape))  
        # if (index == self.selected_layer):  
        #     return x
```

Use the “get_feature” function and delete the part of extracting features to make the process of inference complete and observe the feature maps in each layer.

x:torch.Size([1, 64, 224, 224])
x:torch.Size([1, 64, 224, 224])
x:torch.Size([1, 64, 224, 224])
x:torch.Size([1, 64, 224, 224])
x:torch.Size([1, 64, 112, 112])
x:torch.Size([1, 128, 112, 112])

⋮

x:torch.Size([1, 128, 112, 112])
x:torch.Size([1, 128, 112, 112])
x:torch.Size([1, 128, 112, 112])
x:torch.Size([1, 128, 56, 56])
x:torch.Size([1, 256, 56, 56])

C.Please list the changes of dimension when feeding a image to the VGG-16.

Answer:

[1, 64, 224, 224]
[1, 64, 112, 112]
[1, 128, 112, 112]
[1, 128, 56, 56]
[1, 256, 56, 56]
[1, 256, 28, 28]
[1, 512, 28, 28]
[1, 512, 14, 14]
[1, 512, 7, 7]

Problem 2 [45/100]

1. Please modify the Prob2.ipynb with the following requirements and set the Cross Entropy Loss, Adam Optimizer 0.002 learning rate and betas [0.5,0.999] to train a classifier :
 - A. Design a model with the following structure. (Example 2-4, Page 21)
 - First Conv. layer: Input: RGB, Output Channel 16, second Conv. layer: Output Channel 32, third Conv. layer: Output Channel 64.
 - FC-Layer1: Input: Defined by the third convolutional Layer, Output: 1200
 - FC-Layer2: Input: From FC- Layer1, Output: 600
 - FC-Layer3: Input: From FC- Layer2, Output: equal to your class size [16/40]
 - B. Save the model and name it as 'Prob2.pth' [4/40] (Example 2-4, Page 23)
 - C. Save the optimizer and name it as 'Prob2_1.pth' [6/40] (Example 2-4, Page 23)

Net ()

Layer type	Input channel	Output channel	Filter size	Stride	padding
Conv1	A	B	3	1	2
ReLU					
AvgPool			2	1	1
Conv2	B	C	2	1	1
ReLU					
AvgPool			2	2	1
Conv3	C	D	2	1	1
ReLU					
AvgPool			2	3	1
Linear1	E	G			
ELU					
Linear2	G	F			
ELU					
Linear3	F	G			

Please crop the parts that you modify in Prob2.ipynb and paste to the solution.docx.

Problem 2

- E. Change the dataset to CIFAR10 and the learning rate :0.0002 [5/40] (Example 2-4, Page 26)
- F. Load the 'Prob2.pth and Prob2_1.pth' obtained from C as pretrained model[8/40] (Hint: torch.load function, Page 33)
- G. Train the model on the CIFAR10 dataset[3/40] (Example 2-4, Page 26)
- H. Save the model and name it as 'Prob2_2.pth' [3/40] (Example 2-4, Page 26)

Solution

Net ()

Layer type	Input channel	Output channel	Filter size	Stride	padding
Conv1	A:3	B:16	3	1	2
ReLU					
AvgPool			2	1	1
Conv2	B:16	C:32	2	1	1
ReLU					
AvgPool			2	2	1
Conv3	C:32	D:64	2	1	1
ReLU					
AvgPool			2	3	1
Linear1	E:64*49	G:1200			
ELU					
Linear2	G:1200	F:600			
ELU					
Linear3	F:600	G:10			

Solution 2A

Design a model with the given structure.

CIFAR10: RGB, 3 channel

(Input channel, Output channel, Kernel size, Stride, Padding)

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 16, 3, 1, 2)
        self.relu1 = nn.ReLU()
        self.avgpool1 = nn.AvgPool2d(2, 1, 1)
        self.conv2 = nn.Conv2d(16, 32, 2, 1, 1)
        self.relu2 = nn.ReLU()
        self.avgpool2 = nn.AvgPool2d(2, 2, 1)
        self.conv3 = nn.Conv2d(32, 64, 2, 1, 1)
        self.relu3 = nn.ReLU()
        self.avgpool3 = nn.AvgPool2d(2, 3, 1)
```

Channel × Width × Height

```
self.fc1 = nn.Linear(64*7*7, 1200)
self.elu1 = nn.ELU()
self.fc2 = nn.Linear(1200, 600)
self.elu2 = nn.ELU()
self.fc3 = nn.Linear(600, 10)
```

10 Classes in CIFAR10

```
def forward(self, x):
    batchsize = x.shape[0]
    x = self.avgpool1(self.relu1(self.conv1(x)))
    x = self.avgpool2(self.relu2(self.conv2(x)))
    x = self.avgpool3(self.relu3(self.conv3(x)))
    # print(x.shape)
    x = x.view(batchsize, -1)
    # print(x.shape)
    x = self.elu1(self.fc1(x))
    x = self.elu2(self.fc2(x))
    x = self.fc3(x)
```

Before feeding features into
Fully-Connected layer, we need to straighten
features as
(Batchsize, Channel × Width × Height)

(Hint: Using “print” can help us
observe the dimensions and can
correctly reshape the features)

```
return x
```

Solution 2B & 2C

Save the model and name it as 'Prob2.pth'

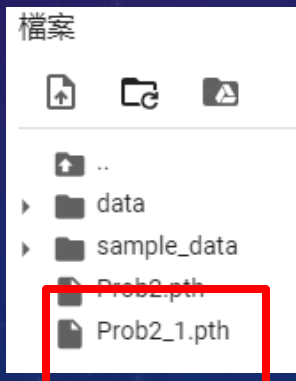
```
Path='Prob2.pth'  
torch.save(net.state_dict(),Path)
```

Define the filename and use “torch.save” to save the model weight file.

Save the optimizer and name it as 'Prob2_1.pth'

```
Path='Prob2_1.pth'  
torch.save(optimizer.state_dict(),Path)
```

Define the filename and use “torch.save” to save the optimizer weight file.



Solution 2E

Change the dataset to CIFAR10 and the learning rate :0.0002

```
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4, shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4, shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.0002, momentum=0.9)
```

Solution 2F

Load the 'Prob2.pth and Prob2_1.pth' obtained from C as pretrained model

```
net = Net() → Initialize the Network
```

```
checkpoint = 'Prob2.pth'
```

```
checkpoint = torch.load(checkpoint) →
```

```
net.load_state_dict(checkpoint) →
```

```
net = net.cuda() → GPU mode
```

```
checkpoint = 'Prob2_1.pth'
```

```
checkpoint = torch.load(checkpoint) →
```

```
optimizer.load_state_dict(checkpoint) →
```

Use “torch.load” to load the model weight file.

Through the line, the net will load the pretrained weights in the model weight file.

Use “torch.load” to load the optimizer weight file.

Through the line, the optimizer will load the pretrained weights in the optimizer weight file.

Solution 2G

Train the model on the CIFAR10 dataset

```
for epoch in range(3): # loop over the dataset multiple times

    running_loss = 0.0
    for i, data in enumerate(trainloader):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data

        inputs = inputs.cuda()
        labels = labels.cuda()

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
        if i % 1000 == 999: # print every 2000 mini-batches

            print("Epoch : {} steps : {} Training Loss : {}".format(epoch + 1, i + 1, running_loss / 1000) )
            running_loss = 0.0

print('Finished Training')
```

→ Covert data to GPU mode

1. Feed the input into the model and get the prediction.
2. Use the defined loss function to calculate the loss between the prediction and the label.
3. Use backward() to compute the gradient and use the optimizer.step() to update the weight

```
Epoch : 1 steps : 1000 Training Loss : 1.0656005302481353
Epoch : 1 steps : 2000 Training Loss : 1.0824362260140479
Epoch : 1 steps : 3000 Training Loss : 1.0971789803281427
Epoch : 1 steps : 4000 Training Loss : 1.0857224909588694
Epoch : 1 steps : 5000 Training Loss : 1.039100713431835
Epoch : 1 steps : 6000 Training Loss : 1.0577989703826607
Epoch : 1 steps : 7000 Training Loss : 1.0318068355247378
```

```
Epoch : 3 steps : 8000 Training Loss : 0.9148140549212694
Epoch : 3 steps : 9000 Training Loss : 0.8839120383523404
Epoch : 3 steps : 10000 Training Loss : 0.905096075758338
Epoch : 3 steps : 11000 Training Loss : 0.9149602136574686
Epoch : 3 steps : 12000 Training Loss : 0.8874567676372827
Finished Training
```

Solution 2H

Save the model and name it as 'Prob2_2.pth'

```
▶ Path='Prob2_2.pth'  
torch.save(net.state_dict(),Path)
```

Define the filename and use “torch.save” to save the second model weight file.

Problem 3 [25/100]

3. The output dimension of the feature map from Conv4 is $64 \times 224 \times 256$, please calculate the dimension of the Input and the feature maps from Conv1, Conv2, Conv3, Conv5

Layer type	Input channel	Output channel	Filter size	Stride
Conv1	3	8	3	1
AvgPool1			4	1
Conv2	8	16	4	2
MaxPool1			2	2
Conv3	16	32	2	1
MaxPool2			2	1
Conv4	32	64	3	2
AvgPool2			3	1
Conv5	64	128	7	1

Solution3

$$\text{Output} = \frac{\text{Input} - \text{kernel size} + 2 \times \text{Paddiing}}{\text{Stride}} + 1$$

$$\text{Input} = (\text{Output} - 1) \times \text{Stride} - 2 \times \text{Paddiing} + \text{kernel size}$$

- Input-> [3,1811,2067]
- Conv1-> [8, 1809, 2065]
- AvgPool1-> [8, 1806, 2062]
- Conv2-> [16, 902, 1030]
- MaxPool1-> [16, 451, 515]
- Conv3-> [32, 450, 514]
- MaxPool2-> [32, 449, 513]
- Conv4-> [64, 224, 256]

$$1811 = (1809 - 1) \times 1 - 2 \times 0 + 3, \quad 2067 = (2065 - 1) \times 1 - 2 \times 0 + 3$$

$$1809 = (1806 - 1) \times 1 - 2 \times 0 + 4, \quad 2065 = (2062 - 1) \times 1 - 2 \times 0 + 4$$

$$1806 = (902 - 1) \times 2 - 2 \times 0 + 4, \quad 2062 = (1030 - 1) \times 2 - 2 \times 0 + 4$$

$$902 = (451 - 1) \times 2 - 2 \times 0 + 2, \quad 1030 = (515 - 1) \times 2 - 2 \times 0 + 2$$

$$451 = (450 - 1) \times 1 - 2 \times 0 + 2, \quad 515 = (514 - 1) \times 1 - 2 \times 0 + 2$$

$$450 = (449 - 1) \times 1 - 2 \times 0 + 2, \quad 514 = (513 - 1) \times 1 - 2 \times 0 + 2$$

$$449 = (224 - 1) \times 2 - 2 \times 0 + 3, \quad 513 = (256 - 1) \times 2 - 2 \times 0 + 3$$

[Return to solution 3](#)

Solution3

- AvgPool2->[64, 222, 254]
- Conv5-> [128, 216, 248]

$$222 = \frac{224 - 3 + 2 \times 0}{1} + 1,$$

$$216 = \frac{222 - 7 + 2 \times 0}{1} + 1,$$

$$254 = \frac{256 - 3 + 2 \times 0}{1} + 1$$

$$248 = \frac{254 - 7 + 2 \times 0}{1} + 1$$