


# *COMPUTER VISION AND ITS APPLICATIONS*

## *ANOMALY DETECTION EXAMPLE AND EXERCISE*

徐繼聖

Gee-Sern Jison Hsu

National Taiwan University of Science  
and Technology






# Recall : f-AnoGAN

- The generator learns the representation of the normal data
  1. Select large volume of normal images  $x$ .
  2. Train the generative model, i.e., generator  $G$  and discriminator  $D$ , on the normal images  $x$ .
  3. Do GAN inversion. Train an encoder  $E$  which can map the image to the latent vector  $z$ .
  4. The encoder  $E$  and the generator  $G$  are used to generate the synthetic normal data  $\hat{x}$ .
  5. Compare test image and its reconstructed image to get anomaly score.

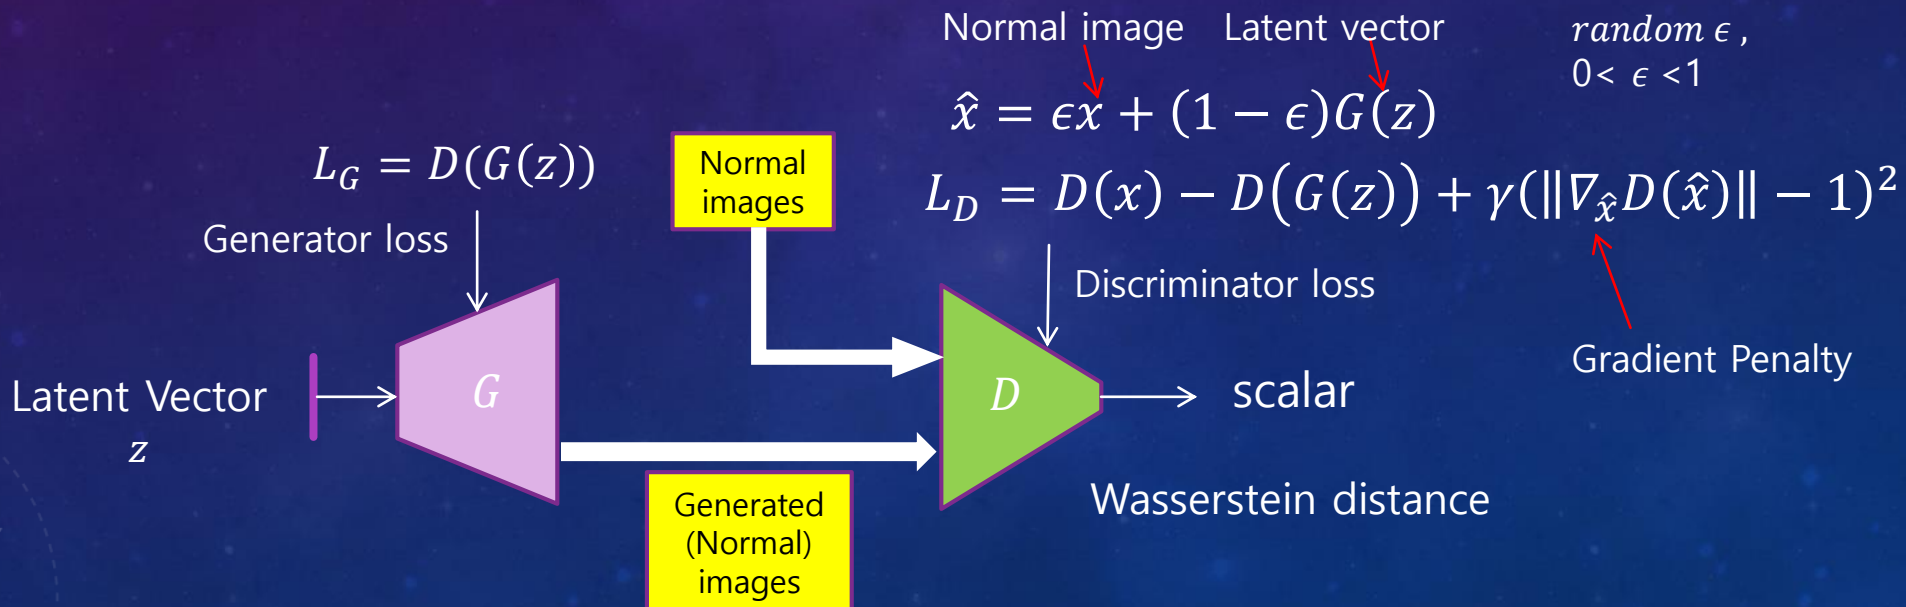
Salient feature: Use GAN inversion to find the mapping of query image to latent space

Anomaly score = feature residual error + image reconstruction error



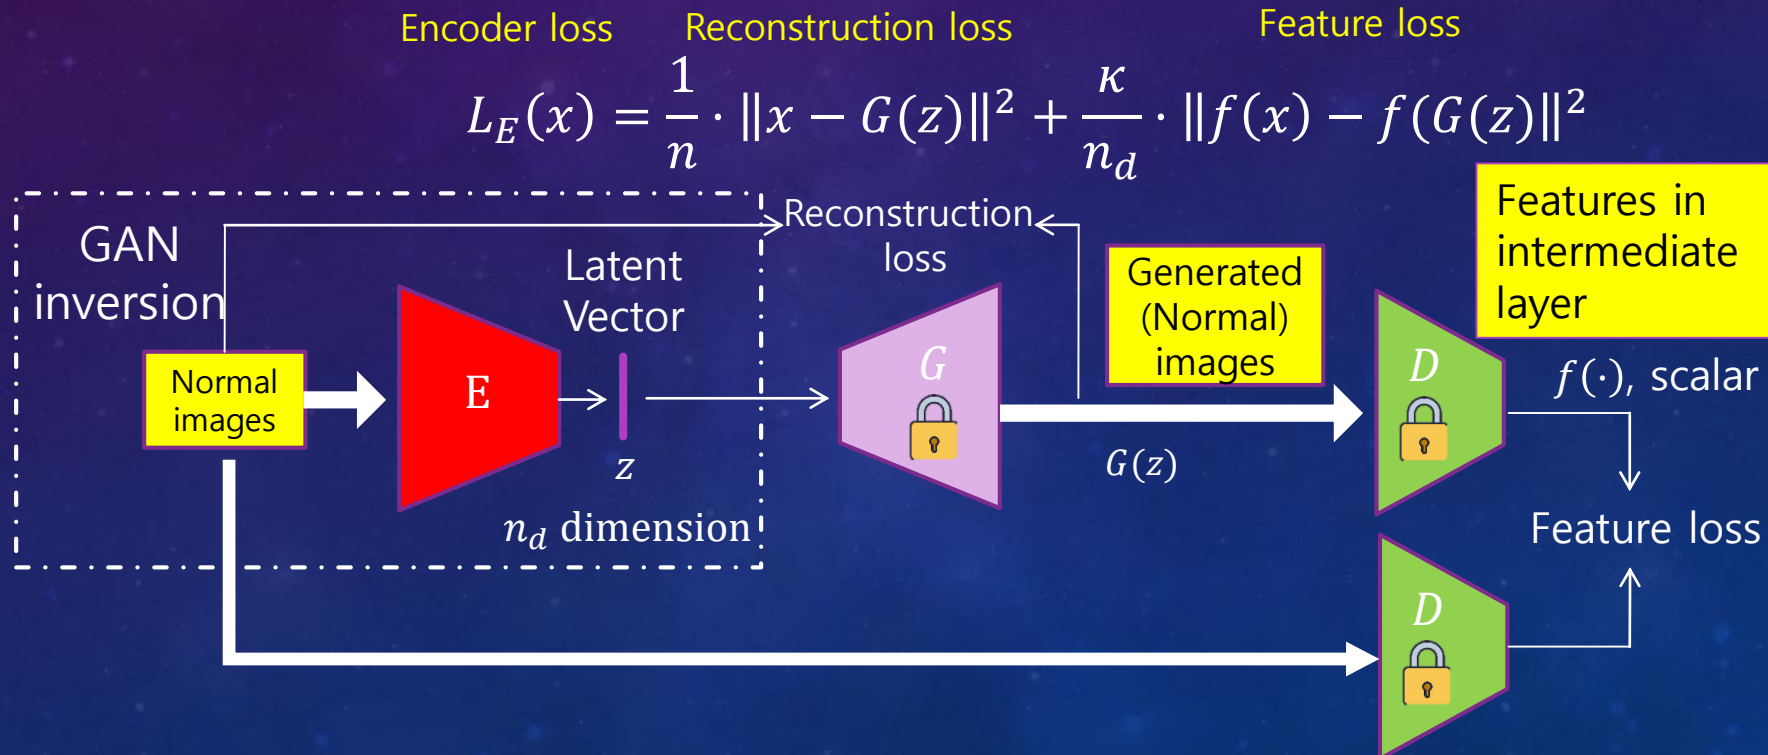
# Recall : f-AnoGAN

- The generator learns the representation of normal image distribution
- The discriminator outputs the Wasserstein distance between generated and real data distribution, not a measure of realness of a given image.
- Use normal image only for the WGAN-GP training
- Latent vector follows the normal distribution.
- Generator & discriminator follow the ResNet network. (In exercise, DCGAN are used instead.)



# Recall : f-AnoGAN


- Learn the mapping of image to a vector in latent space where the generator is trained from.
- Parameters in generator and discriminator are fixed.
- The generated image shall closely resemble to the input normal image
- The discriminator outputs the features in the intermediate layer, not Wasserstein distance nor a measure of image realness.





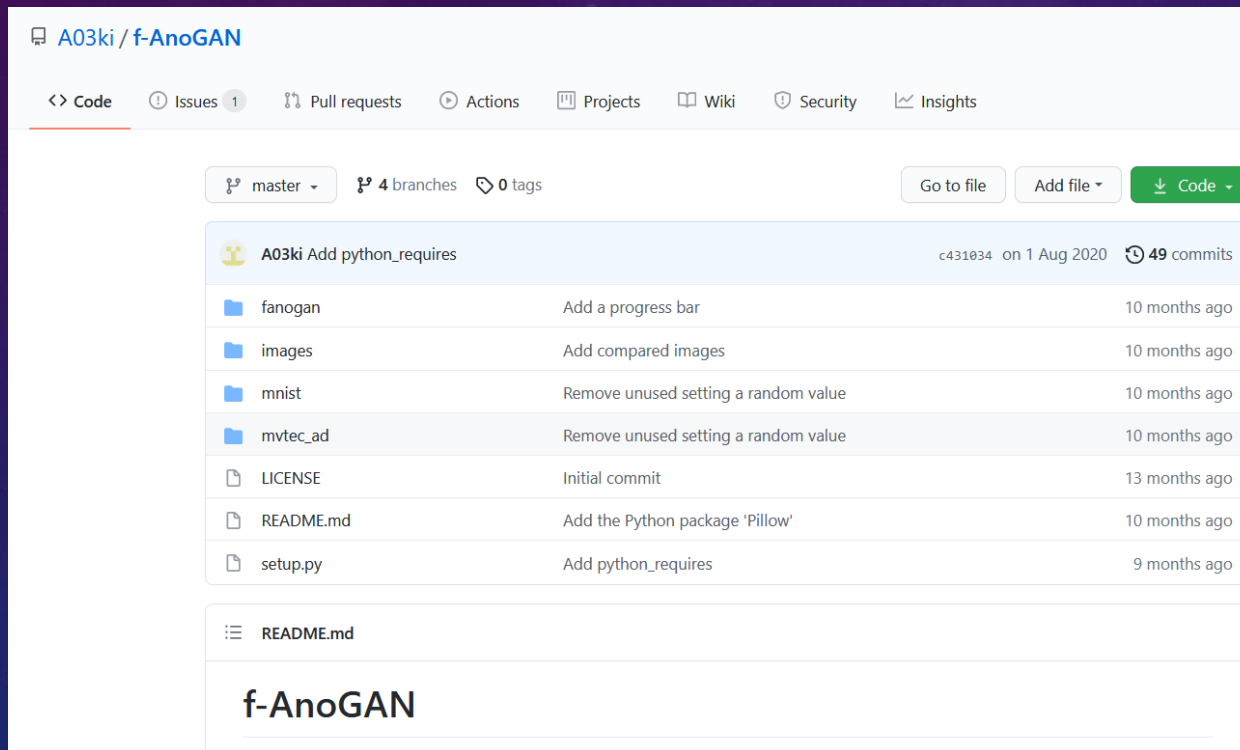
# Recall : f-AnoGAN

- Anomaly score (AS) for image-level detection is the sum of reconstruction loss and feature loss of query image.
- If the query image is a normal image, the AS shall be as smaller as possible.
- The AS of anomalous query image shall be significantly different from the AS of normal images.
- For the pixel-level anomaly localization, we use the absolute value of pixel-wise residual,  $\hat{A}_R = |x - G(z)|$



# Exercise\_MNIST : f-AnoGAN

- Please visit <https://github.com/A03ki/f-AnoGAN> .
- Scroll down to the bottom of page and enter "f-AnoGAN\_MNIST.ipynb"



# Exercise\_MNIST : f-AnoGAN - Pre-requisites

## Pre-requisites :

Follow the instruction to download the codes of f-AnoGAN and setup the environments.

### Step: 0

Please push "Open in playground" and run below in order.

1.

```
[23] !git clone https://github.com/A03ki/f-AnoGAN.git
```

```
Cloning into 'f-AnoGAN'...
remote: Enumerating objects: 194, done.
remote: Counting objects: 100% (194/194), done.
remote: Compressing objects: 100% (125/125), done.
remote: Total 194 (delta 108), reused 139 (delta 67), pack-reused 0
Receiving objects: 100% (194/194), 220.66 KiB | 3.87 MiB/s, done.
Resolving deltas: 100% (108/108), done.
```

2.

```
[24] %cd f-AnoGAN
```

```
/content/f-AnoGAN/mnist/f-AnoGAN/f-AnoGAN
```

3.

```
[25] !python setup.py install
```

4.

```
[4] %cd mnist
```

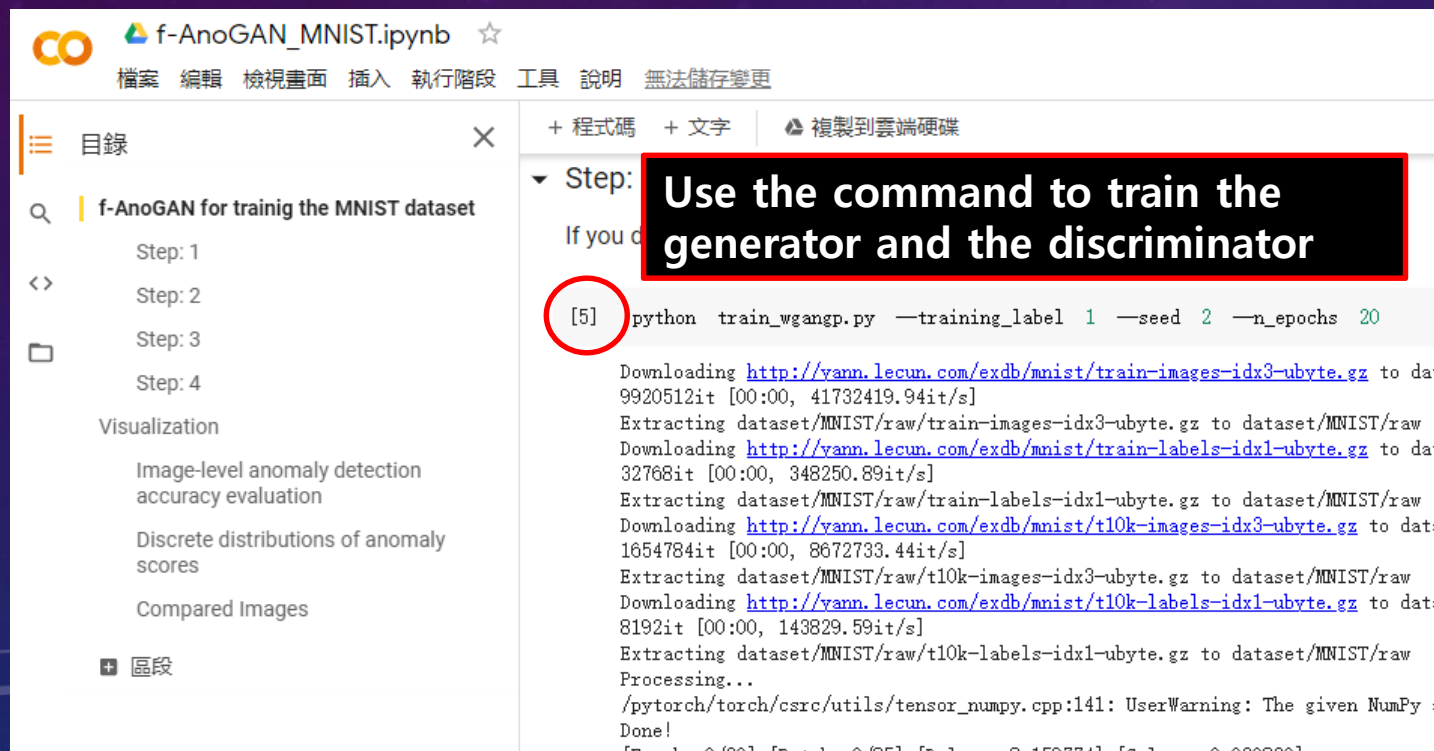
4.

```
/content/f-AnoGAN/mnist
```

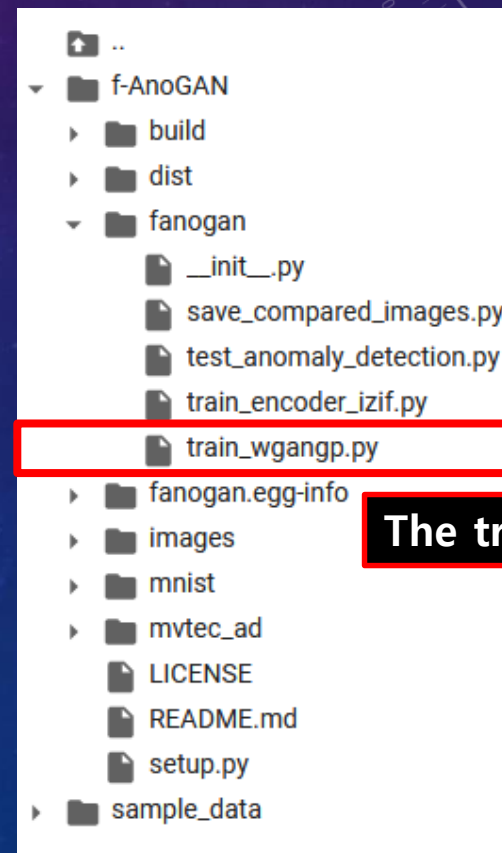
# Exercise\_MNIST : f-AnoGAN – Train WGANGP (1/5)

## Step 1. Train WGANGP model

Use the command "python train\_wgangp.py ..." to train the generator and the discriminator. All the codes can be visualized in "f-AnaGAN/fanongan/train\_wgangp.py"



The screenshot shows a Jupyter Notebook titled "f-AnoGAN\_MNIST.ipynb". The left sidebar contains a table of contents with sections: "f-AnoGAN for training the MNIST dataset", "Step: 1", "Step: 2", "Step: 3", "Step: 4", "Visualization", "Image-level anomaly detection accuracy evaluation", "Discrete distributions of anomaly scores", "Compared Images", and "區段". The main cell displays a code block with the command: `[5] python train_wgangp.py --training_label 1 --seed 2 --n_epochs 20`. A red circle highlights the cell number [5]. A red box with white text is overlaid on the code, stating: "Use the command to train the generator and the discriminator". Below the command, there is a large block of text showing the progress of downloading and extracting MNIST dataset files.





# Exercise\_MNIST : f-AnoGAN – Train WGANGP (2/5)

f-AnoGAN\_MNIST.ipynb ☆

檔案 編輯 檢視畫面 插入 執行階段 工具 說明 無法儲存變更

目錄

Step: 1

training\_label == 1, which means we select the class "1" as the normal class. (Other numbers are abnormal)

Step: 2

Step: 3

Step: 4

Visualization

Image-level anomaly detection accuracy evaluation

Discrete distributions of anomaly scores

Compared Images

+ 區段

```
[5] !python train_wgangp.py --training_label 1 --seed 2 --n_epochs 20
```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to dat  
9920512it [00:00, 41732419.94it/s]  
Extracting dataset/MNIST/raw/train-images-idx3-ubyte.gz to dataset/MNIST/raw  
Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to dat  
32768it [00:00, 348250.89it/s]  
Extracting dataset/MNIST/raw/train-labels-idx1-ubyte.gz to dataset/MNIST/raw  
Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to dat  
1654784it [00:00, 8672733.44it/s]  
Extracting dataset/MNIST/raw/t10k-images-idx3-ubyte.gz to dataset/MNIST/raw  
Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to dat  
8192it [00:00, 143829.59it/s]  
Extracting dataset/MNIST/raw/t10k-labels-idx1-ubyte.gz to dataset/MNIST/raw  
Processing...  
/pytorch/torch/csrc/utils/tensor\_numpy.cpp:141: UserWarning: The given NumPy a  
Done!  
[Epoch: 0/20] [Batch: 0/85] [D\_loss: 8.159774] [G\_loss: 0.020820]

# Exercise\_MNIST : f-AnoGAN – Train WGANGP (3/5)

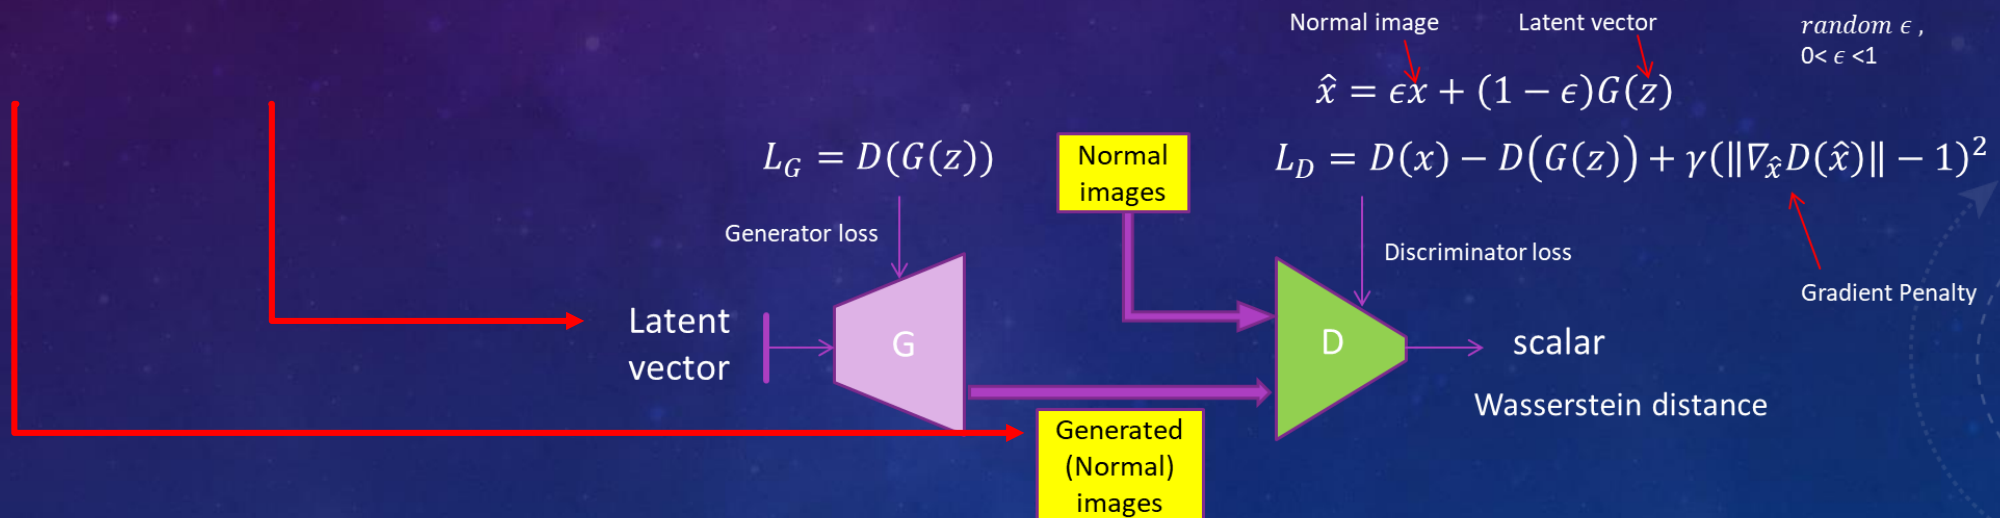
```
optimizer_D.zero_grad()

# Sample noise as generator input
z = torch.randn(imgs.shape[0], opt.latent_dim, device=device)

# Generate a batch of images
fake_imgs = generator(z)
```

During training the generator and the discriminator, we use the random noise  $z$  to generate the synthesized normal data.

The generator  $G$  aims to synthesize the normal data from the random noise  $z$





# Exercise\_MNIST : f-AnoGAN – Train WGAN GP (4/5)

Train the discriminator.

```
# Real images
real_validity = discriminator(real_imgs)
# Fake images
fake_validity = discriminator(fake_imgs.detach())
# Gradient penalty
gradient_penalty = compute_gradient_penalty(discriminator,
                                           real_imgs.data,
                                           fake_imgs.data,
                                           device)
# Adversarial loss
d_loss = (-torch.mean(real_validity) + torch.mean(fake_validity)
          + lambda_gp * gradient_penalty)

d_loss.backward()
optimizer_D.step()

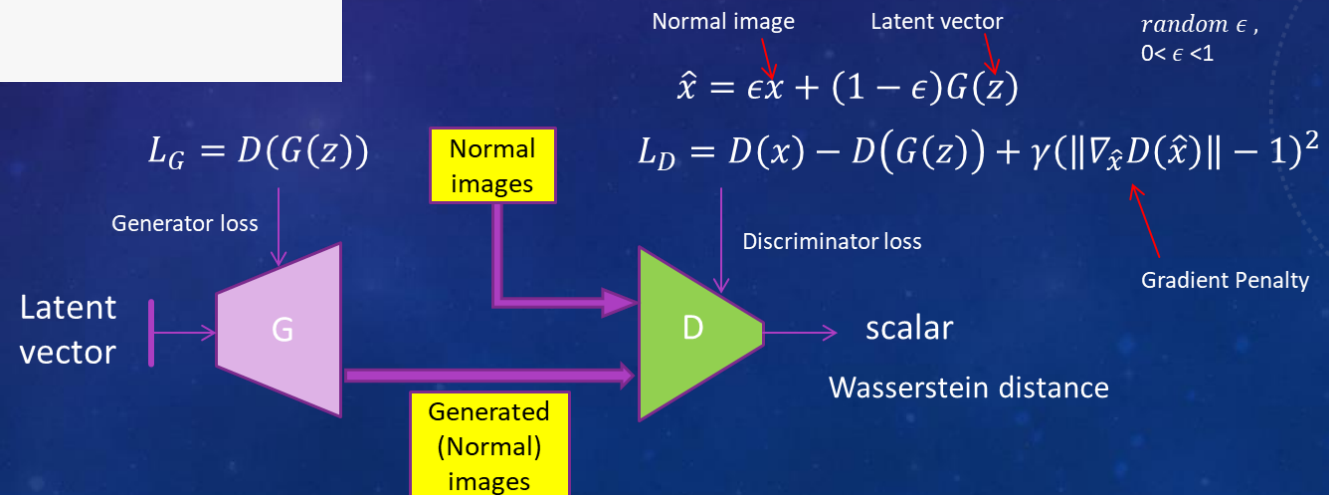
optimizer_G.zero_grad()
```

$D(x)$  denotes the output of the real image.

$D(G(x))$  denotes the output of the generated image.

The loss function, WGAN-GP

Gradient penalty is  $\nabla_{\hat{x}} D(\hat{x})$



# Exercise\_MNIST : f-AnoGAN – Train WGANGP (5/5)

Train the generator

```
# Train the generator and output log every n_critic steps
if i % opt.n_critic == 0:

    # -----
    # Train Generator
    # -----

    # Generate a batch of images
    fake_imgs = generator(z)
    # Loss measures generator's ability to fool the discriminator
    # Train on fake images
    fake_validity = discriminator(fake_imgs)
    g_loss = -torch.mean(fake_validity)

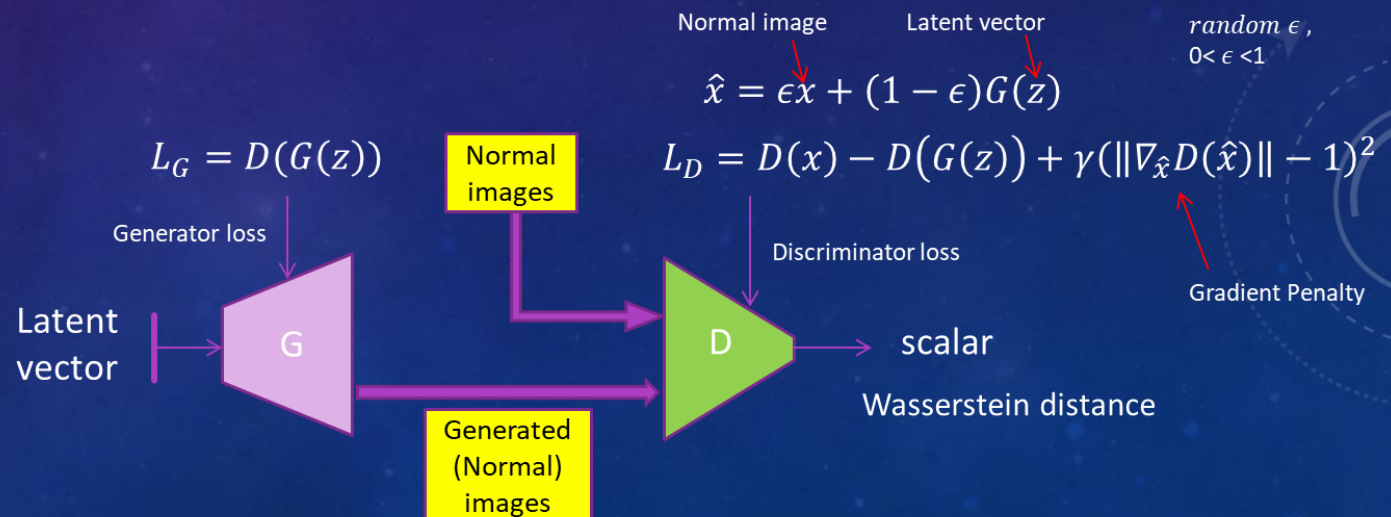
    g_loss.backward()
    optimizer_G.step()

    print(f"[Epoch {epoch:{padding_epoch}}/{opt.n_epochs}] "
          f"[Batch {i:{padding_i}}/{len(data_loader)}] "
          f"[D loss: {d_loss.item():3f}] "
          f"[G loss: {g_loss.item():3f}]")

    if batches_done % opt.sample_interval == 0:
        save_image(fake_imgs.data[:25],
                    f"results/images/{batches_done:06}.png",
                    nrow=5, normalize=True)

    batches_done += opt.n_critic
```

$G$  tries to fool  $D \rightarrow L_G = D(G(z))$



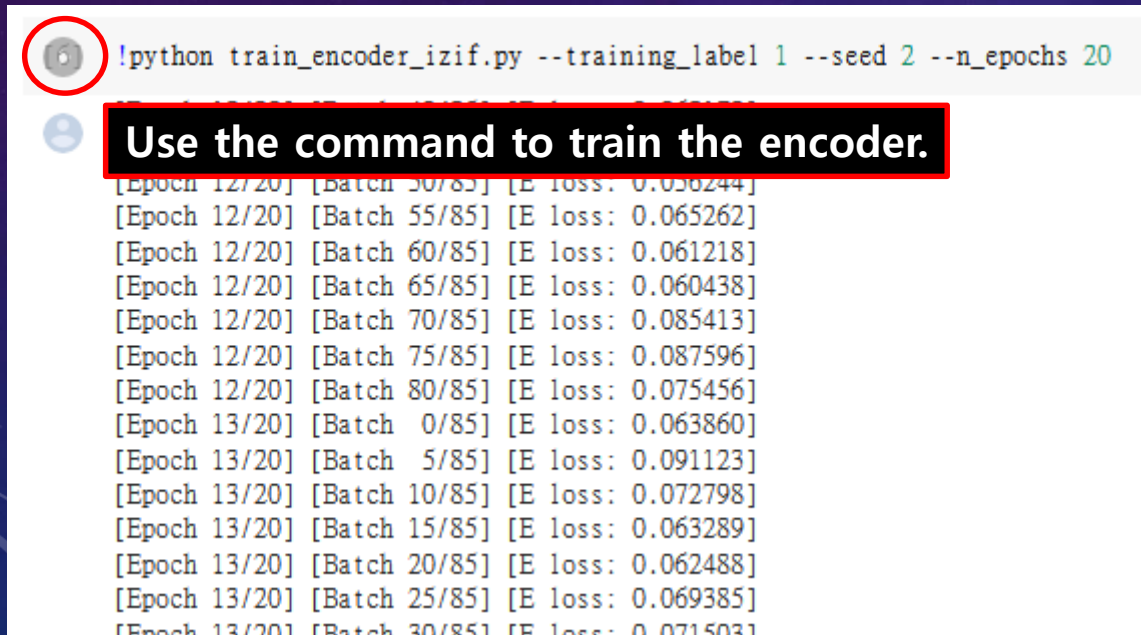


# Exercise\_MNIST: f-AnoGAN – Train Encoder (1/2)

## Step 2. Train encoder model

The encoder aims to learn how to map the real image to a latent vector, which can be used to generate the images by the generator.

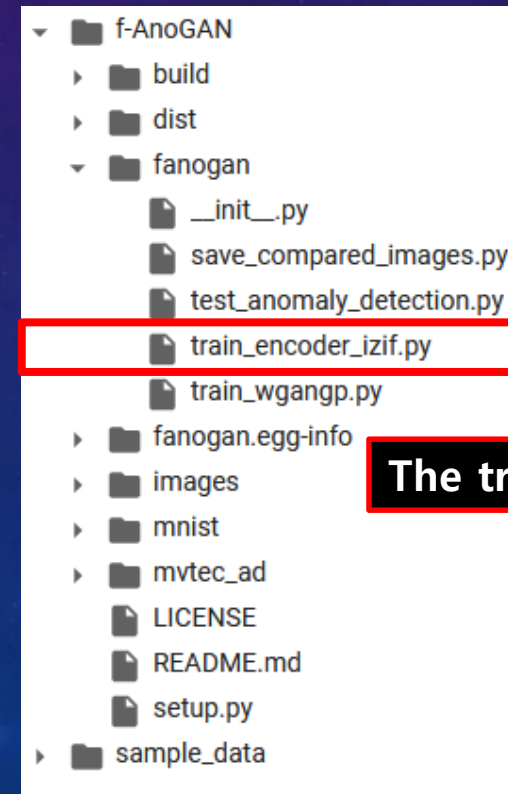
Use the command "python train\_encoder\_izif.py ..." to train the encoder. The training codes can be visualized in "f-AnaGAN/fanongan/train\_encoder\_izif.py"



```
!python train_encoder_izif.py --training_label 1 --seed 2 --n_epochs 20
```

**Use the command to train the encoder.**

```
[Epoch 12/20] [Batch 50/85] [E loss: 0.056244]  
[Epoch 12/20] [Batch 55/85] [E loss: 0.065262]  
[Epoch 12/20] [Batch 60/85] [E loss: 0.061218]  
[Epoch 12/20] [Batch 65/85] [E loss: 0.060438]  
[Epoch 12/20] [Batch 70/85] [E loss: 0.085413]  
[Epoch 12/20] [Batch 75/85] [E loss: 0.087596]  
[Epoch 12/20] [Batch 80/85] [E loss: 0.075456]  
[Epoch 13/20] [Batch 0/85] [E loss: 0.063860]  
[Epoch 13/20] [Batch 5/85] [E loss: 0.091123]  
[Epoch 13/20] [Batch 10/85] [E loss: 0.072798]  
[Epoch 13/20] [Batch 15/85] [E loss: 0.063289]  
[Epoch 13/20] [Batch 20/85] [E loss: 0.062488]  
[Epoch 13/20] [Batch 25/85] [E loss: 0.069385]  
[Epoch 13/20] [Batch 30/85] [E loss: 0.071503]
```



**The training code.**

# Exercise\_MNIST: f-AnoGAN – Train Encoder (2/2)

We use two objective functions to train the encoder:

- 1) Reconstruction loss
- 2) Feature loss.

```
# Real features
real_features = discriminator.forward_features(real_imgs)
# Fake features
fake_features = discriminator.forward_features(fake_imgs)

# izif architecture
loss_imgs = criterion(fake_imgs, real_imgs)
loss_features = criterion(fake_features, real_features)
e_loss = loss_imgs + kappa * loss_features

e_loss.backward()
optimizer_E.step()
```

Reconstruction loss

Feature loss



# Exercise\_MNIST: f-AnoGAN – Inference (1/2)

## Step 3. Inference

Use the command "test\_anomaly\_detection.py ..." to test the performance. The inference codes can be visualized in "f-AnoGAN/fanongan/ test\_anomaly\_detection.py"

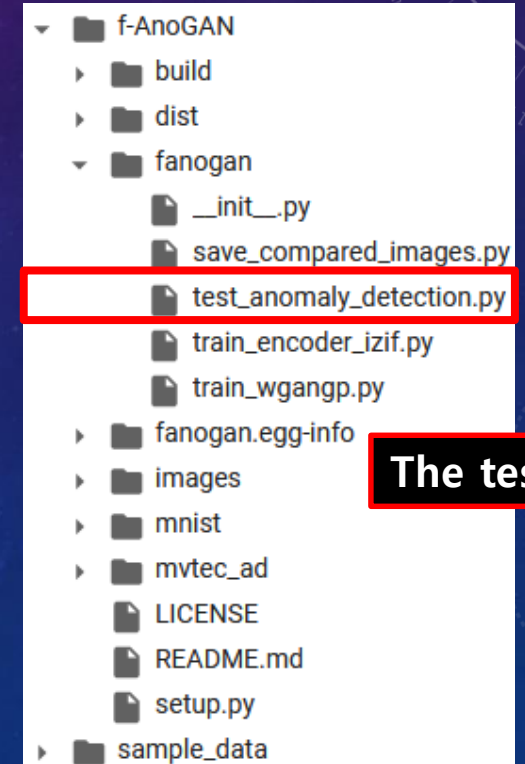
### Step: 3

If you don't run command `test_anomaly_detection.py` yet, please run below after Step: 2.

[7] !python **Use the command to test the model.**

100% 64607/64607 [02:28<00:00, 436.15it/s]

Scores for anomaly detection are saved under `f-AnoGAN/mnist/results`.



**The test code.**

# Exercise\_MNIST: f-AnoGAN – Inference (2/2)

```
6 def test_anomaly_detection(opt, generator, discriminator, encoder,  
7                             dataloader, device, kappa=1.0):  
8     generator.load_state_dict(torch.load("results/generator"))  
9     discriminator.load_state_dict(torch.load("results/discriminator"))  
10    encoder.load_state_dict(torch.load("results/encoder"))  
11  
12    generator.to(device).eval()  
13    discriminator.to(device).eval()  
14    encoder.to(device).eval()  
15  
16    criterion = nn.MSELoss()  
17  
18    with open("results/score.csv", "w") as f:  
19        f.write("label,img_distance,anomaly_score,z_distance\n")  
20
```

Load the trained weights

```
21 for (img, label) in tqdm(dataloader):  
22  
23     real_img = img.to(device)  
24  
25     real_z = encoder(real_img)  
26     fake_img = generator(real_z)  
27     fake_z = encoder(fake_img)  
28  
29     real_feature = discriminator.forward_features(real_img)  
30     fake_feature = discriminator.forward_features(fake_img)  
31  
32     # Scores for anomaly detection  
33     img_distance = criterion(fake_img, real_img)  
34     loss_feature = criterion(fake_feature, real_feature)  
35     anomaly_score = img_distance + kappa * loss_feature  
36  
37     z_distance = criterion(fake_z, real_z)
```

Calculate the error metrics  
(distances and anomaly scores)



# Exercise\_MNIST: f-AnoGAN – Visualize (1/3)

**Step 4. Obtain the test results and analyze**, the codes can be visualized in "f-AnaGAN/fanongan/save\_compared\_images.py"

## ▼ Step: 4

If you don't run command `save_compared_images.py` yet, please run below after Step: 2.

[8] `!python save_compared_images.py --seed 4 --n_iters 0 --n_grid_lines 10`

1.

Compared images are saved under `f-AnoGAN/mnist/results/images_diff`.

## ▼ Visualization

Please run below after Step: 1~3.

[9] `import matplotlib.pyplot as plt`  
`import numpy as np`  
2. `import pandas as pd`  
`from sklearn.metrics import roc_curve, precision_recall_curve, auc`

[10] `df = pd.read_csv("results/score.csv")`  
`df` 3.

4. "1" is normal class

	label	img_distance	anomaly_score	z_distance
0	1	0.072017	0.129134	0.008652
1	1	0.033213	0.036116	0.003864
2	1	0.034159	0.042097	0.028874
3	1	0.030439	0.035003	0.003410
4	1	0.067693	0.095846	0.012227
...	...	...	...	...
64602	2	0.521851	0.815865	0.779736
64603	3	0.482058	0.622129	0.382944
64604	4	0.436425	0.515557	0.188704
64605	5	0.376133	0.488694	0.144208
64606	6	0.638045	0.930252	0.156074
64607 rows × 4 columns				

others are abnormal

# Exercise\_MNIST: f-AnoGAN – Visualize (2/3)


Calculate the results.

```
[26] trainig_label = 1
      labels = np.where(df["label"].values == trainig_label, 0, 1)
5.    anomaly_score = df["anomaly_score"].values
      img_distance = df["img_distance"].values
      z_distance = df["z_distance"].values
```

```
▶ fpr, tpr, _ = roc_curve(labels, img_distance)
  precision, recall, _ = precision_recall_curve(labels, img_distance)
6.    roc_auc = auc(fpr, tpr)
      pr_auc = auc(recall, precision)
```

## Image-level anomaly detection accuracy evaluation


```
[13] plt.plot(fpr, tpr, label=f"AUC = {roc_auc:3f}")
      plt.plot([0, 1], [0, 1], linestyle="--")
7.    plt.title("ROC-AUC")
      plt.xlabel("False Positive Rate")
      plt.ylabel("True Positive Rate")
      plt.legend()
      plt.show()
```



# Exercise\_MNIST: f-AnoGAN – Visualize (3/3)

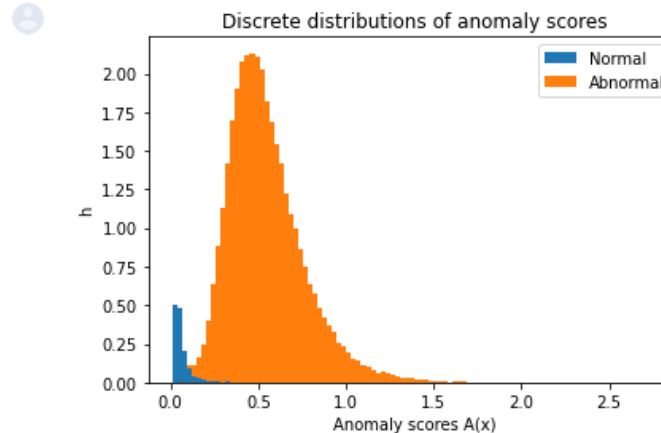
Draw the ROC curve, data distribution and the real/generated sample test.

```
[14] plt.plot(recall, precision, label=f"PR = {pr_auc:3f}")  
plt.title("PR-AUC")  
8. plt.xlabel("Recall")  
plt.ylabel("Precision")  
plt.legend()  
plt.show()
```




## Discrete distributions of anomaly scores

```
[15] plt.hist([anomaly_score[labels == 0], anomaly_score[labels == 1]],  
            bins=100, density=True, stacked=True,  
            label=["Normal", "Abnormal"])  
9. plt.title("Discrete distributions of anomaly scores")  
plt.xlabel("Anomaly scores A(x)")  
plt.ylabel("h")  
plt.legend()  
plt.show()
```



```
[17] from PIL import Image
```


```
10. Image.open("results/images_diff/000010.png")
```





# Homework : GANomaly


- GANomaly detects anomalies with a novel encoder-decoder-encoder structure.
- The detection criterion is similar to f-AnoGAN.
- The distance of latent vector  $z$  and  $\hat{z}$  is defined as the anomaly score.



$$\mathcal{L}_{enc} = \|z - \hat{z}\|$$

$$\mathcal{L}_{con} = \|x - \hat{x}\|_1$$

$$\mathcal{L}_{adv} = \|f(x) - f(\hat{x})\|_2$$



# Homework : GANomaly

The code use Cifar10 dataset as default.

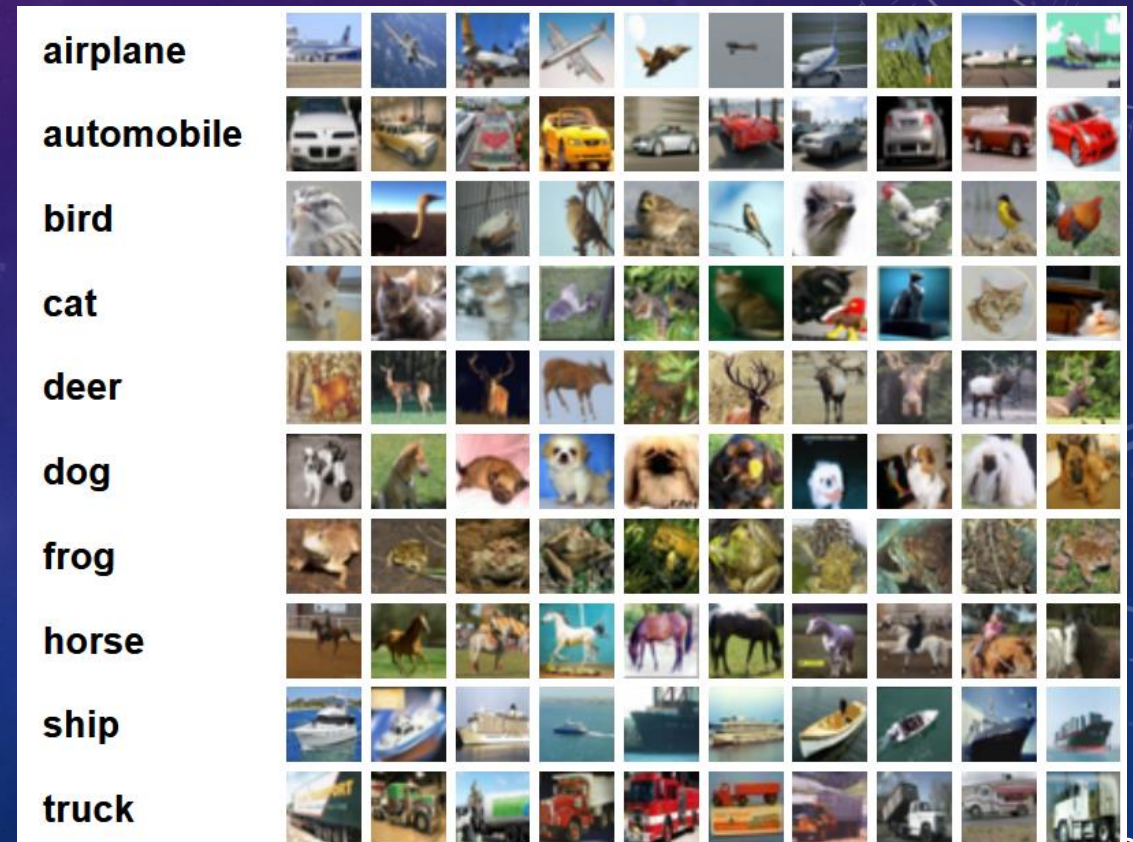
The dataset consists of 60k 32x32 color images in 10 classes, with 6k per class. (5k training / 1k test)

GANomaly define 'car' as the abnormal class (All 6000 images)

The other classes follow 5k training 1k test protocol.

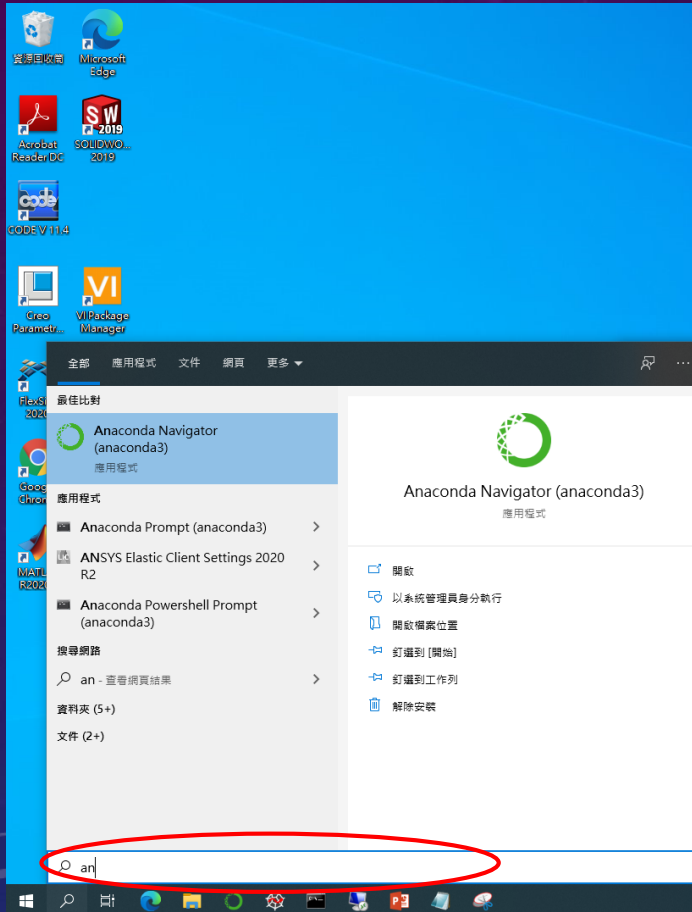
Training : 45,000 normal images

Test : 6,000 abnormal + 9,000 normal images

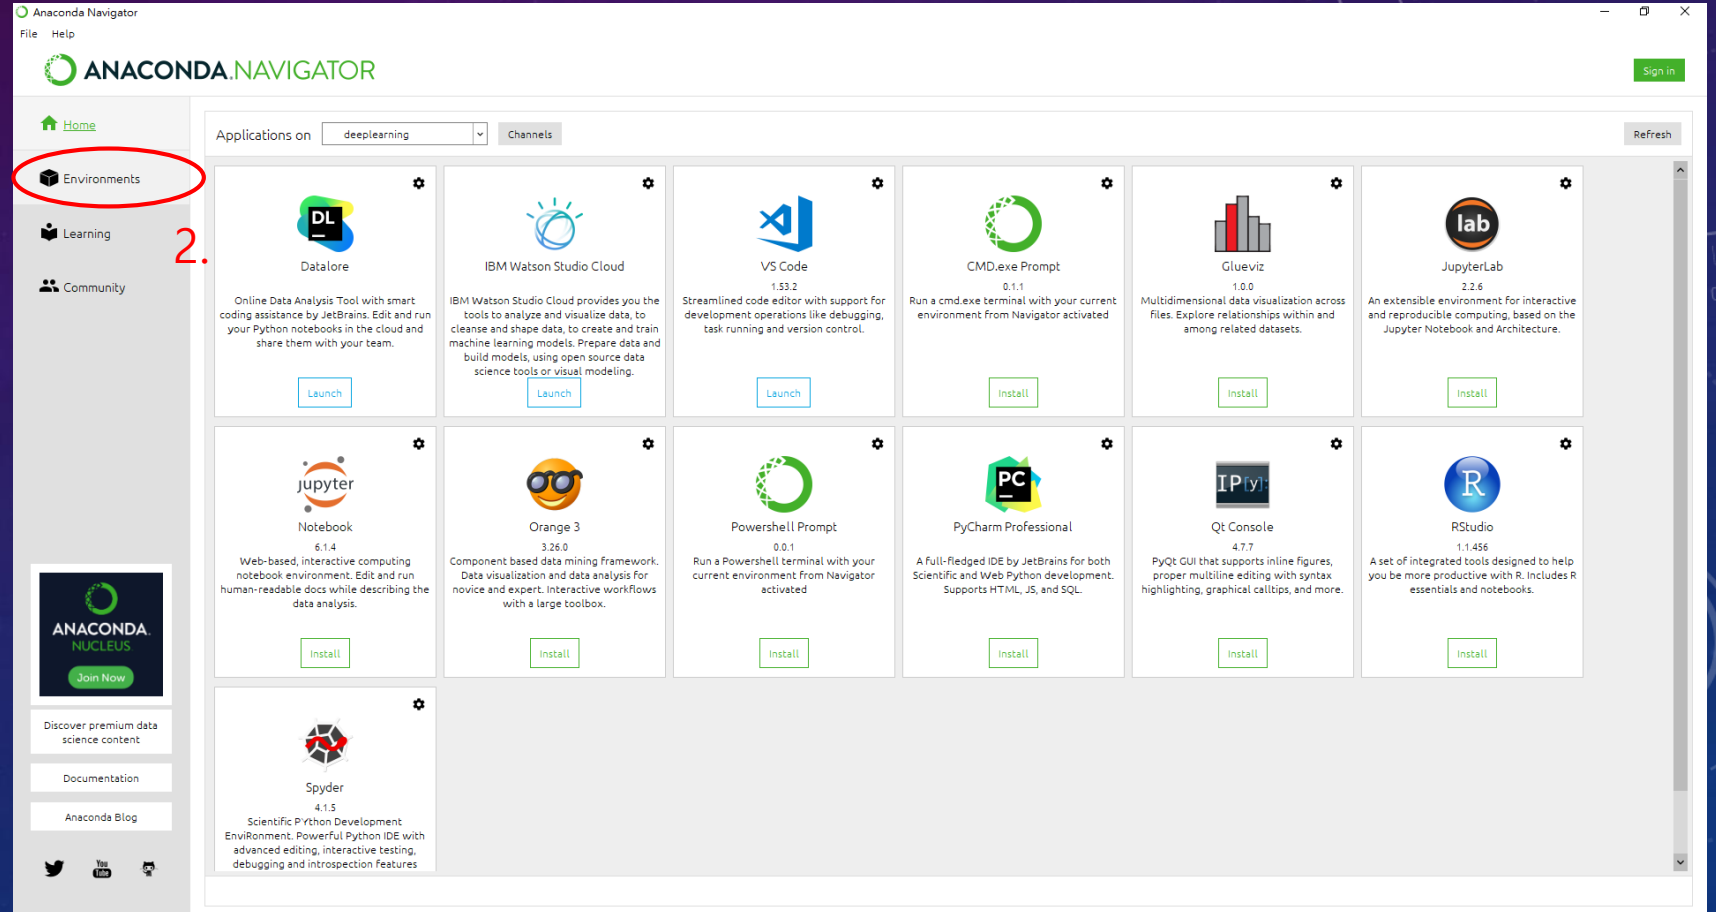


# Homework : GANomaly

0. Please download "ganomaly-master\_modified.zip" from moodle.

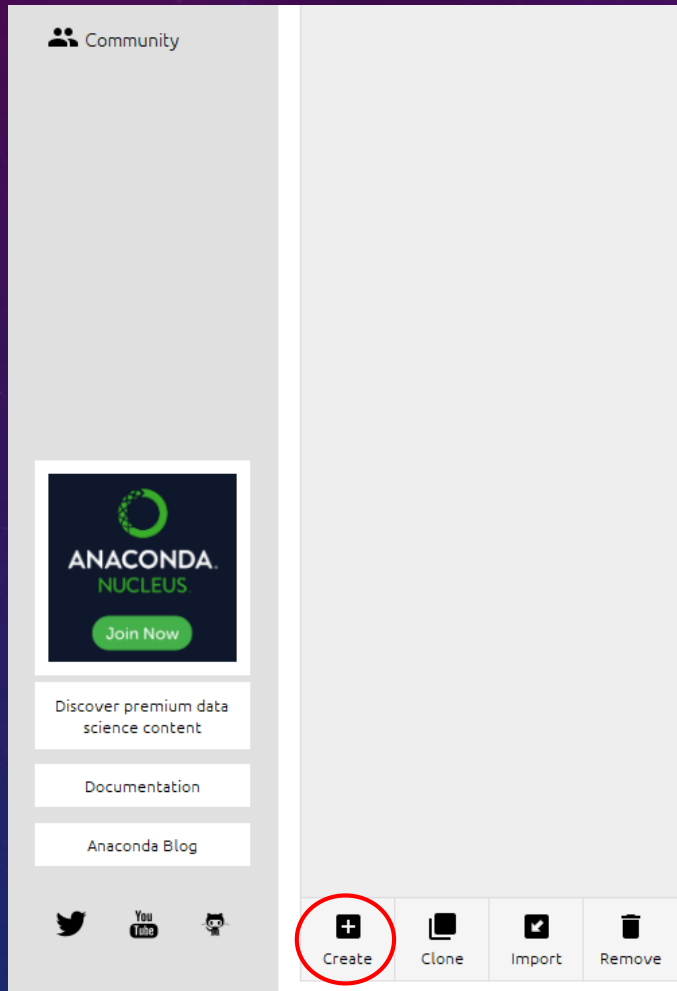


1. Enter "Anaconda"

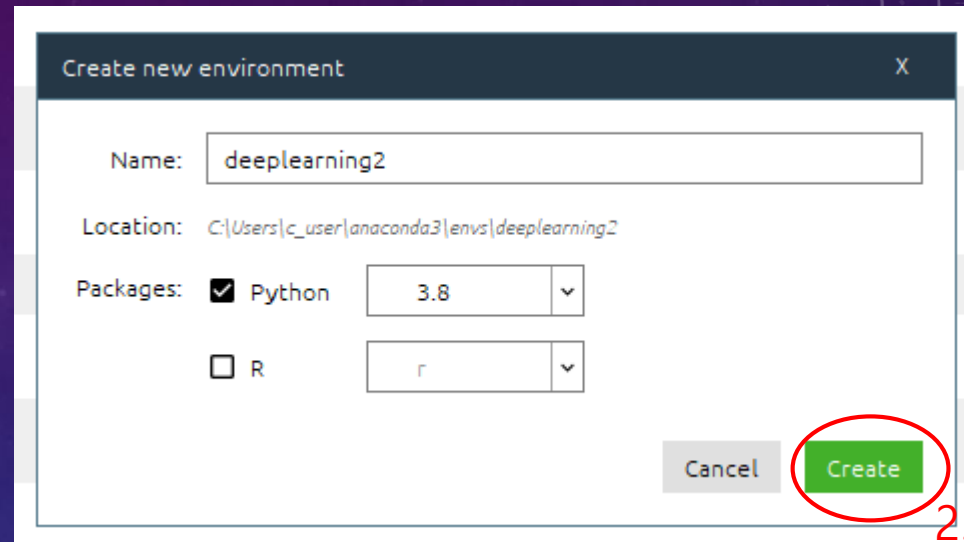





# Homework : GANomaly



1.



2.



3.

# Homework : GANomaly

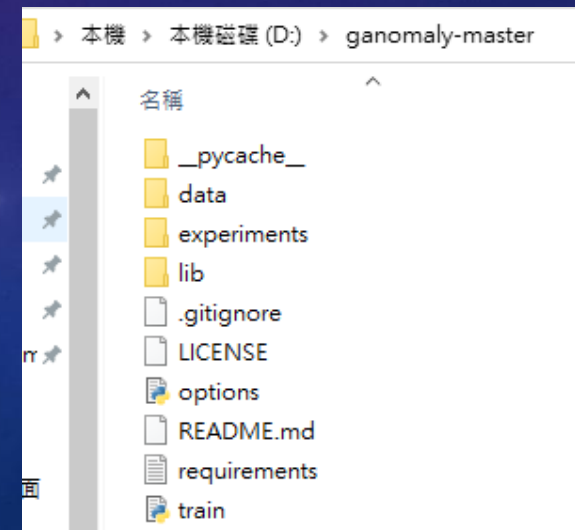

```
C:\Windows\system32\cmd.exe
```

```
(deeplearning) C:\Users\c_user>D:
```

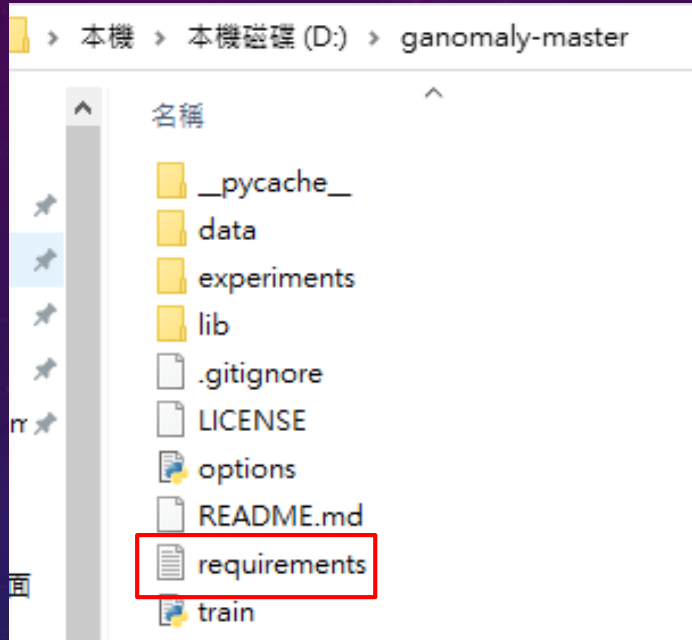
```
(deeplearning) D:\>cd ganomaly-master
```

```
(deeplearning) D:\ganomaly-master>
```

Change  
directory



# Homework : GANomaly



Please remove these items due to version issues.

"torch" and "scipy"

1.

```
asn1crypto==0.24.0
certifi==2019.6.16
cffi==1.12.3
chardet==3.0.4
cryptography==2.7
cyclr==0.10.0
idna==2.8
joblib==0.13.2
kiwisolver==1.1.0
matplotlib==3.1.0
numpy==1.16.4
olefile==0.46
Pillow==7.1.0
pycparser==2.19
pyOpenSSL==19.0.0
pyparsing==2.4.0
PySocks==1.7.0
python-dateutil==2.8.0
pytz==2019.1
pyzmq==18.0.2
requests==2.22.0
scikit-learn==0.21.2
scipy==1.3.0
six==1.12.0
torch==1.2.0
torchfile==0.1.0
torchvision==0.4
tornado==6.0.3
tqdm==4.33.0
urllib3==1.25.3
visdom==0.1.8.8
websocket-client==0.56.0
```

## 2. After 1., we install the required packages

2-1.

```
(deeplearning) D:\ganomaly-master>pip install -r requirements.txt
```

2-2.

```
(deeplearning) D:\ganomaly-master>
```

```
conda install pytorch torchvision torchaudio cudatoolkit=10.2 -c pytorch
```

2-3.

```
(deeplearning) D:\ganomaly-master>pip install tqdm scipy tensorboardX tensorflow
```



# Homework : GANomaly

After steps above, training is ready.

```
(deeplearning) D:\ganomaly-master>python train.py
```

Please open another terminal and start Tensorboard. (After training is started.)

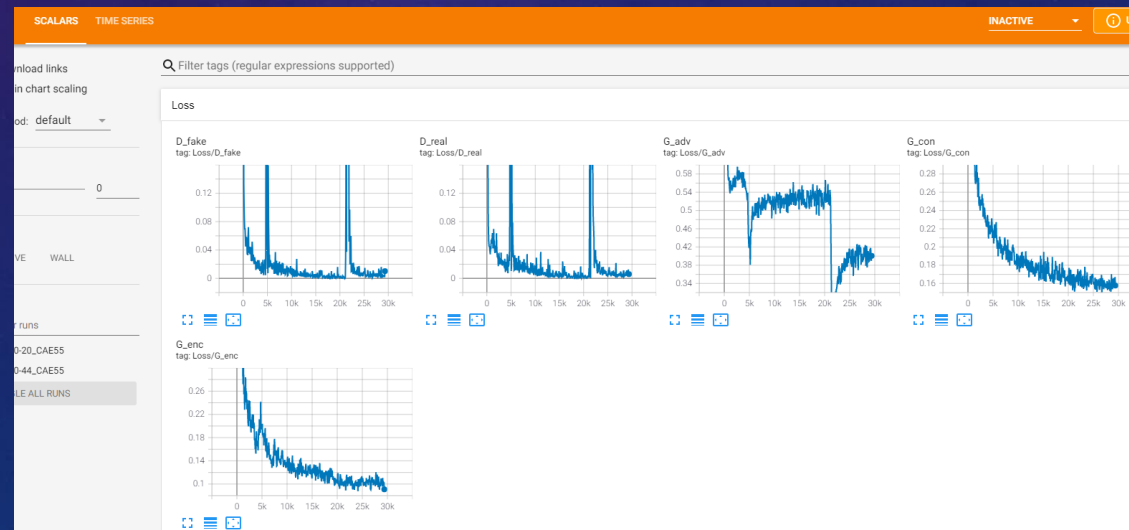
```
(deeplearning) D:\ganomaly-master>cd runs
```

```
(deeplearning) D:\ganomaly-master\runs>tensorboard --logdir . --port 6006
```

```
Serving TensorBoard on localhost; to expose to the network, use a proxy or pass --bind_all  
TensorBoard 2.5.0 at http://localhost:6006/ (Press CTRL+C to quit)
```

\*Tensorboard is a tool that collects the loss values for further analysis

Please copy the address and paste it to your browser.



# Homework 4-2: GANomaly

Please show your training loss curve, and discuss the peaks of the graphs.  
(Hint : The relation between G and D)



# Homework : GANomaly

Please show the images of epoch 1, 10 or more epochs. And discuss what may be happening? Is there any relation between loss and fake images?

Fake\_image

Fake\_image

step 0 Mon May 10 2021 16:07:12 GMT+0800 (台北標準時間)



Fake\_image

step 9

Mon May 10 2021 16:26:09 GMT+0800 (台北標準時間)






# Homework : GANomaly

Please compare your results with the pre-trained model.

The pertained model is already in "output/ganomaly/cifar10/train/weights/"



```
(pytorch_high) D:\ganomaly-master>python train.py --load_weights --phase test
Files already downloaded and verified
Files already downloaded and verified
>> Training model GANomaly.
    Loaded weights.
OrderedDict([('Avg Run Time (ms/batch)', 9.974679946899414), ('roc', 0.692968462962963)])
```

Phase test for pretrain model

```
(pytorch_high) D:\ganomaly-master>python train.py --load_weights --phase valid
Files already downloaded and verified
Files already downloaded and verified
>> Training model GANomaly.
    Loaded weights.
```


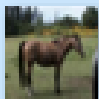

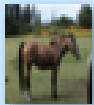
Phase valid for your model

< Your results >

# Homework : GANomaly

Given a pre-trained model shown in last page, please compare the anomaly scores and the generated image between the pre-trained model and the model your trained, and write down the observation in your .docx file.

1. Please compare the generated image from your model and the pretrained model
2. Please compare the anomaly scores between these two models.
3. Write down what have you been observed.

	Anomaly Score	
Pre-trained	 Anomaly_test_0.101.png	 Normal_test_0.032.png
Yours	 Anomaly_valid_0.075.png	 Normal_valid_0.030.png